

# Networks

## Congestion Control

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Congestion Control

Sections 3.6-3.7

# TCP Congestion Control

- Underlying network might be overwhelmed by aggregate traffic load
  - Packet delays due to queues in routers
  - Packet loss due to buffer overflow at routers
- Not the same as flow control
- Rather than making situation worse with retransmissions, slow down transmission when the network is congested

# TCP Congestion Control

- Congestion window ( $cwnd$ ) limits how much unACKed data can be sent
- Maximum segment size (MSS)
- Additive increase: increase  $cwnd$  by 1 MSS every RTT until loss detected
- Multiplicative decrease: cut  $cwnd$  in half after loss detected


# Maximum Segment Size (MSS)

- Maximum transmission unit (MTU) is the maximum frame for a given link layer (e.g., 1500 bytes for Ethernet and PPP)
- Maximum segment size is typically sender's MTU minus 40 bytes for TCP/IP headers
- Path MTU (RFC 1191) is largest frame supported by all links in a given path

# TCP Congestion Control


- AIMD: probe for usable bandwidth until loss occurs

AIMD saw tooth  
behavior: probing  
for bandwidth



# TCP Slow Start


- Increase transmission rate exponentially until first loss event
  - Initially,  $cwnd = 1$  MSS
  - Double  $cwnd$  every RTT (i.e., increment for each ACK received)





# TCP Congestion Avoidance

- Switch from exponential increase to linear increase when `cwnd` gets to  $\frac{1}{2}$  value before timeout
- Implemented with `ssthresh` variable set to  $\frac{1}{2}$  of `cwnd` just before loss event



# TCP Congestion Control

- Loss indicated by timeout
  - `cwnd` set to 1 MSS
  - Window grows exponentially (slow start) to `ssthresh`, then linearly (congestion avoidance)
- Loss indicated by triple duplicate ACKs
  - TCP RENO cuts `cwnd` in half, then grows linearly from there (TCP fast recovery)
  - TCP TAHOE reacts as if timeout-based loss

# TCP Flow Control

- Avoid overflowing receiver with too much data sent too fast
- Not the same as congestion control
- Receiver advertises free buffer space `rwnd` to sender in window size field of TCP header

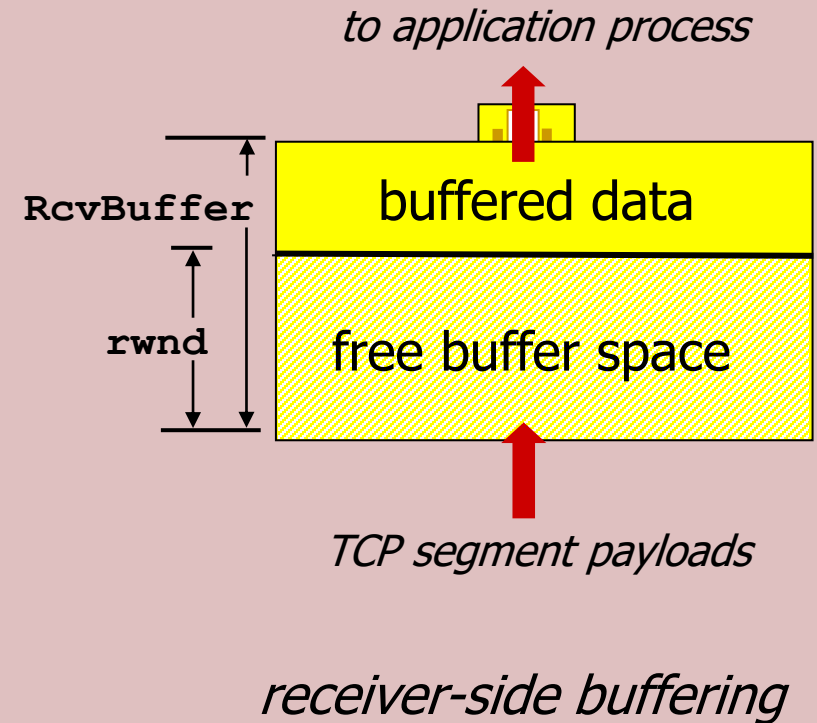
# TCP Flow Control

- Receiver buffer

`RcvBuffer`


- Set via socket options (default is typically 4KB)
- Some OSes auto adjust `RcvBuffer`

- Sender will send with  $W = \min(\text{cwnd}, \text{rwnd})$



# TCP Fairness

- Goal: If  $K$  TCP connections share same bottleneck link of bandwidth  $R$ , then each connection *should* have an average rate of  $R/K$



# TCP (Un-)Fairness

- Multimedia applications often use UDP
  - Typically can tolerate *some* packet loss
  - Avoid throttling due to congestion control
- Multiple parallel TCP connections
  - Browser implementations do this
  - Example: Suppose 9 connections share link with capacity  $R$ 
    - App 1 requests one connection ( $R/10$  per app)
    - App 2 requests 10 connections ( $R/2$  for app 2)
    - Single connection apps experience  $R/20$

# QUIC Review

- Quick UDP Internet Connections
- Transport-like services implemented over UDP at the application layer
- Developed by Google initially as an experimental protocol
- Undergoing IETF standardization

# QUIC

- Initially targeted improved transport performance between Chrome browser and Google services
- Over 1/3 of Google egress traffic (over 5% of internet traffic)
- Akamai CDN deployed in 2016



# QUIC vs. TCP

- Many transport services overlap (congestion control, reliability, etc.)
- Multiple streams in single QUIC connection
- QUIC uses 64-bit connection IDs rather than 4-tuples that simplify migrating to different addresses and ports

# Thank You!

# Networks

Connectionless Transport - UDP

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Connectionless Transport - UDP

## Section 3.3


# User Datagram Protocol

- Minimal best-effort transport protocol
  - Datagrams may be lost (unreliable)
  - Datagrams may arrive out of order
- Connectionless
  - No connection setup or teardown overhead
  - Each datagram handled independently

# User Datagram Protocol

- No congestion or flow control
- Applications must implement desired services (e.g., reliability, flow control)
  - Streaming multimedia
  - DNS
  - QUIC (not really an application)

# UDP Datagram Format



What is the largest possible UDP packet size?



# UDP Checksum

- Also known as *Internet checksum* or *TCP checksum*
- Detect bit errors in datagrams
- Sender: set checksum header field to be checked by the receiver
- Receiver: verify checksum for received datagram

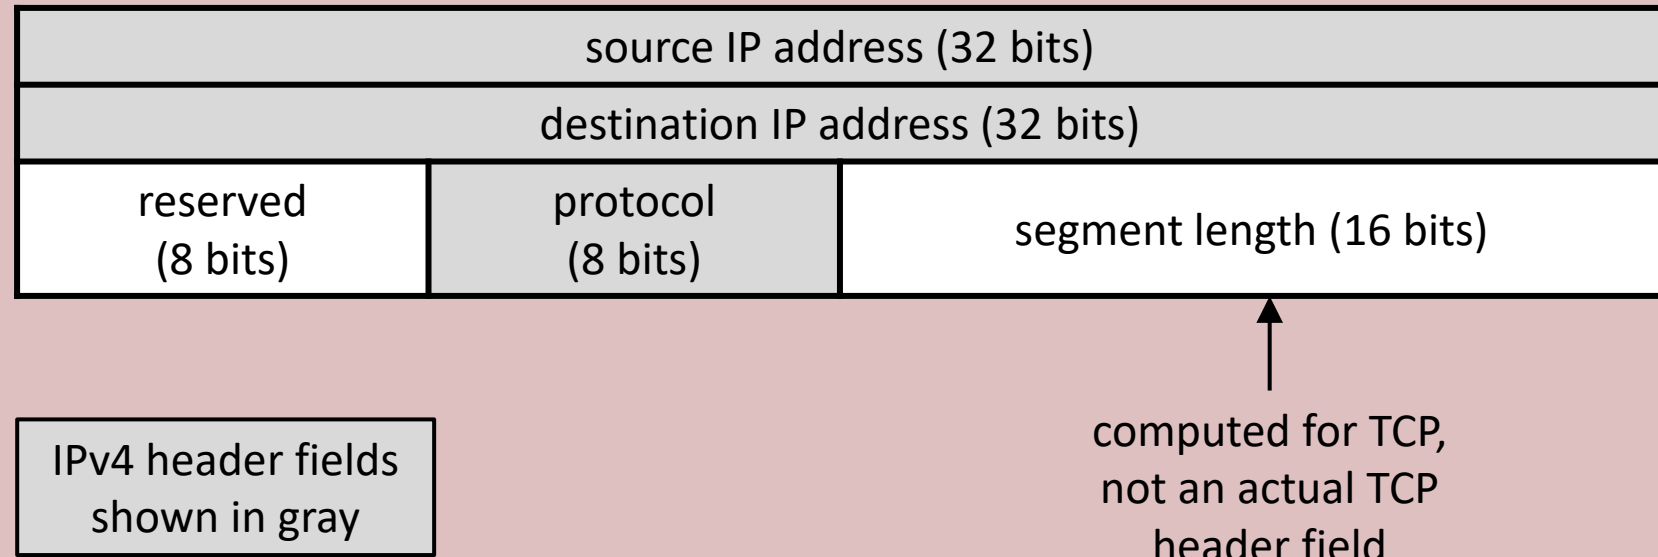
# UDP Checksum

- One's complement of the one's complement sum of the following data treated as a sequence of 16-bit integers
  - Header (checksum field treated as all 0s for calculation)
  - Payload
  - Pad byte (if odd number of payload bytes)
  - Pseudo-header


# UDP Pseudo-Header

- Used for checksum computation *but not actually transmitted*
- Mostly IP header fields
  - 12 bytes for IPv4
  - 40 bytes for IPv6
- Protocol layering violation (why?)

# UDP Pseudo-Header (IPv4)



# UDP Pseudo-Header



# UDP Checksum

Example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

# Thank You!

# Networks

Connectionless Transport - UDP



Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Connectionless Transport - UDP

## Section 3.3


# User Datagram Protocol

- Minimal best-effort transport protocol
  - Datagrams may be lost (unreliable)
  - Datagrams may arrive out of order
- Connectionless
  - No connection setup or teardown overhead
  - Each datagram handled independently

# User Datagram Protocol

- No congestion or flow control
- Applications must implement desired services (e.g., reliability, flow control)
  - Streaming multimedia
  - DNS
  - QUIC (not really an application)

# UDP Datagram Format



What is the largest possible UDP packet size?

# UDP Checksum

- Also known as *Internet checksum* or *TCP checksum*
- Detect bit errors in datagrams
- Sender: set checksum header field to be checked by the receiver
- Receiver: verify checksum for received datagram

# UDP Checksum


- One's complement of the one's complement sum of the following data treated as a sequence of 16-bit integers
  - Header (checksum field treated as all 0s for calculation)
  - Payload
  - Pad byte (if odd number of payload bytes)
  - Pseudo-header

# UDP Pseudo-Header

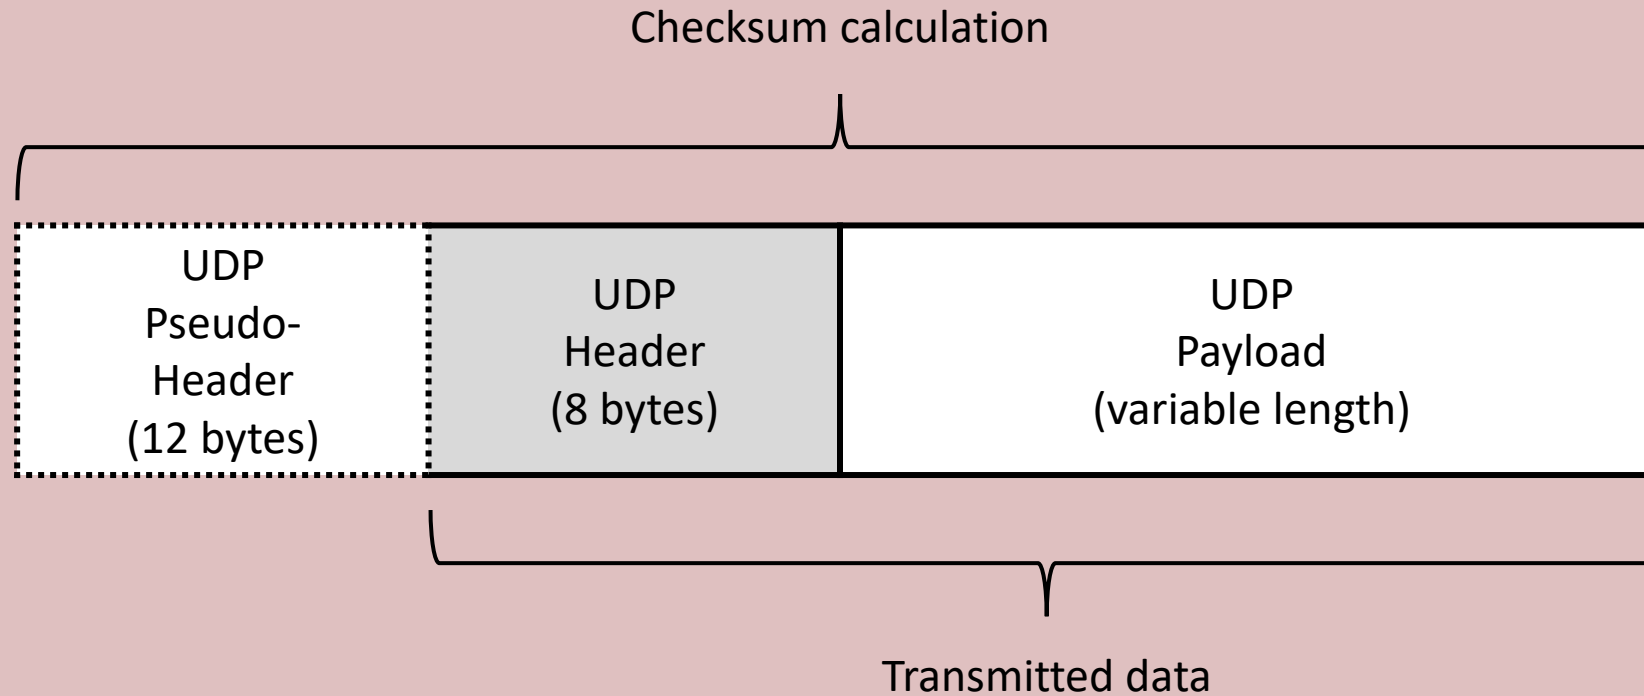
- Used for checksum computation *but not actually transmitted*
- Mostly IP header fields
  - 12 bytes for IPv4
  - 40 bytes for IPv6
- Protocol layering violation (why?)



# UDP Pseudo-Header (IPv4)



# UDP Pseudo-Header



# UDP Checksum

Example: add two 16-bit integers

		1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
		1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																	
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
<hr/>																	
sum		1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum		0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

# Thank You!

# Networks

Transport-Layer Services

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Transport-Layer Services

## Section 3.1

# Transport Layer

- Runs only on the end systems
- Provides logical process-to-process communication
  - Sender breaks application messages into UDP datagrams or TCP segments
  - Receiver reassembles for application
- Built on top of network layer services (host-to-host communication)



# Inter-Office Mail Analogy

Company split between NY and LA offices

- Hosts = offices
- Processes = employees
- Application messages = company memos
- Transport layer = local mailroom that delivers individual envelopes to desks
- Network-layer = U.S. postal service

# Transport Layer Protocols

- UDP: essentially datagram headers added to IP packets
- TCP: reliable, connection-oriented over IP
- QUIC: reliable, connection-oriented protocol over UDP (i.e., technically application layer)

# Transport Layer Services

Service	UDP	TCP	QUIC
Reliable delivery		✓	✓
In-order delivery		✓	✓
Congestion control		✓	✓
Flow control		✓	✓
Connection-oriented		✓	✓
Delay guarantees			
Bandwidth guarantees			

# Thank You!

# Networks

## Multiplexing and Demultiplexing

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Multiplexing and Demultiplexing

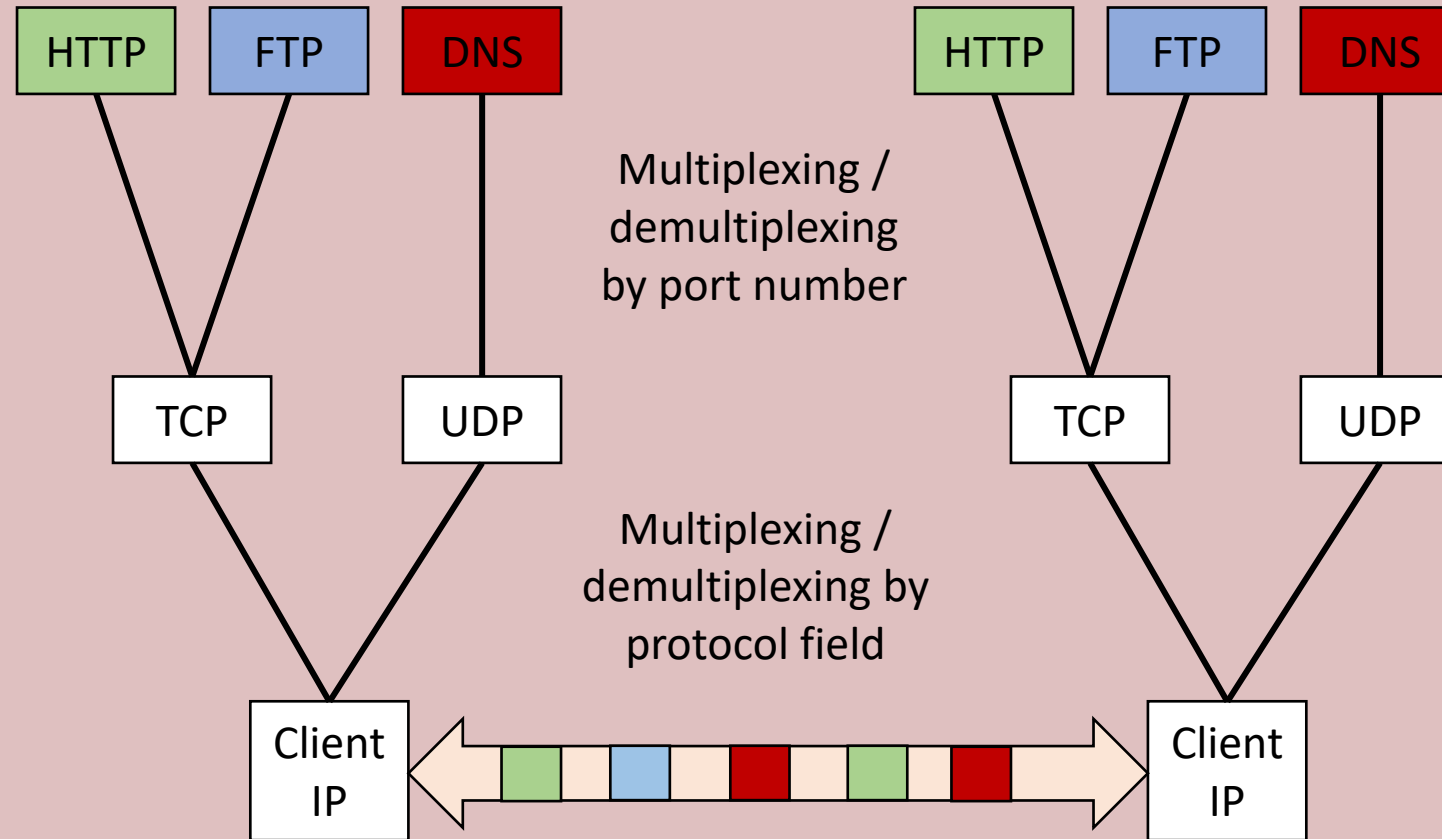
## Section 3.2

# Protocol Multiplexing and Demultiplexing

- Multiplexing at sender
  - Handling messages from multiple sockets
  - Handling datagrams/segments from TCP or UDP
- Demultiplexing at receiver
  - Delivering packets to either UDP or TCP based on protocol number in IP header
  - Delivering datagrams/segments to correct socket based on port number in UDP/TCP header



# Protocol Multiplexing and Demultiplexing



# Thank You!

# Networks

Transmission Control Protocol

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Transmission Control Protocol

## Section 3.5

# Transmission Control Protocol (TCP)

- Originally published in RFC 793 in 1981
- Variant of Go-Back-N but also has SACK option for Selective Repeat
- Connection-oriented
- Reliable, in-order byte stream service
- Full duplex between two endpoints
- Congestion control
- Flow control

# TCP Segment Header

source port (16 bits)			destination port (16 bits)		
sequence number (32 bits)					
acknowledgement number (32 bits)					
length (4 bits)	unused (4 bits)	flags (8 bits)		window size (16 bits)	
checksum (16 bits)			urgent pointer (16 bits)		
options (variable length)					

# TCP Segment Header

- Length: header length in 32-bit words (options are variable length), basic length is 20 bytes without any options
- Urgent pointer: rarely used method for marking special data
- Window size: number of bytes receiver is willing to accept (more later)
- Checksum: identical to UDP checksum, but over more fields
- Flags: mostly used for connection management



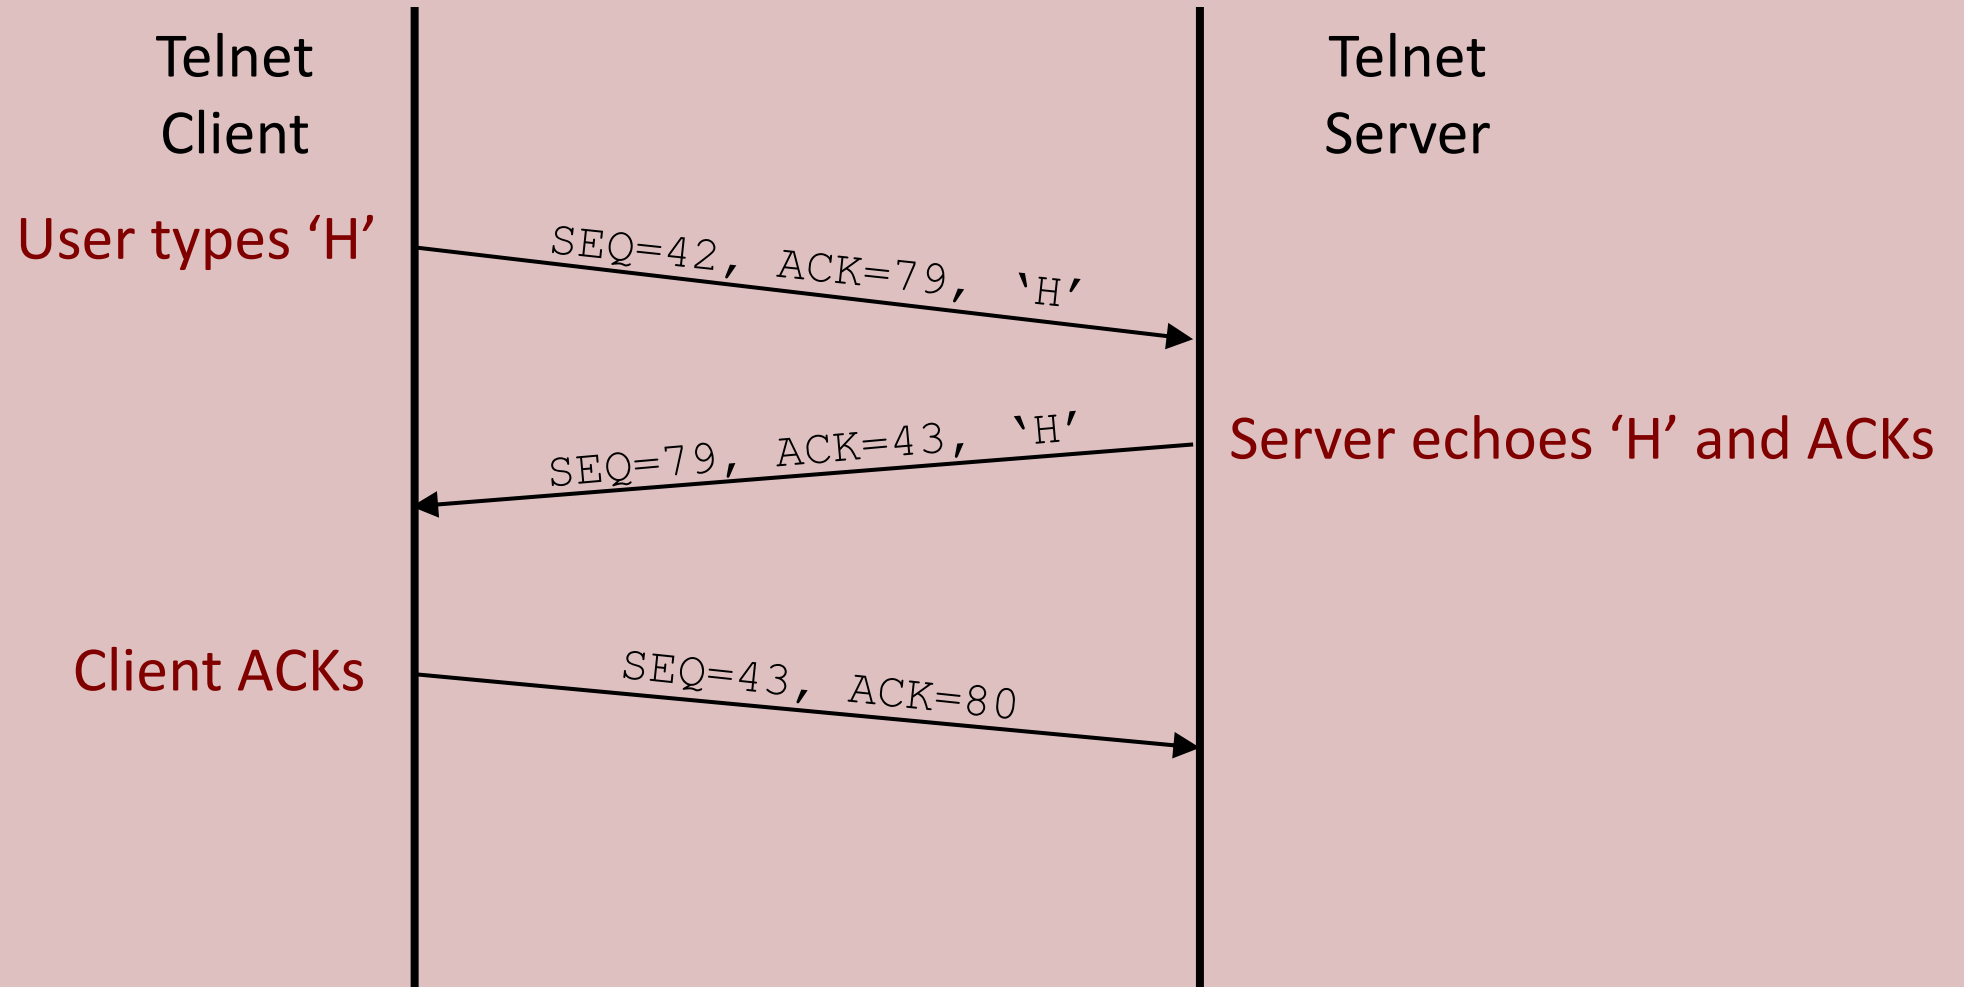
# TCP Sequence Number

- 32-bit unsigned integer in TCP segment header
- Represents byte number in byte stream that first byte of segment represents
- Wraps back around to 0

# TCP Acknowledgement Number

- Next sequence number expected to be received (cumulative ACK)
- Alternatively, sequence number of last successfully received byte of data plus 1
- Only valid if ACK flag is set

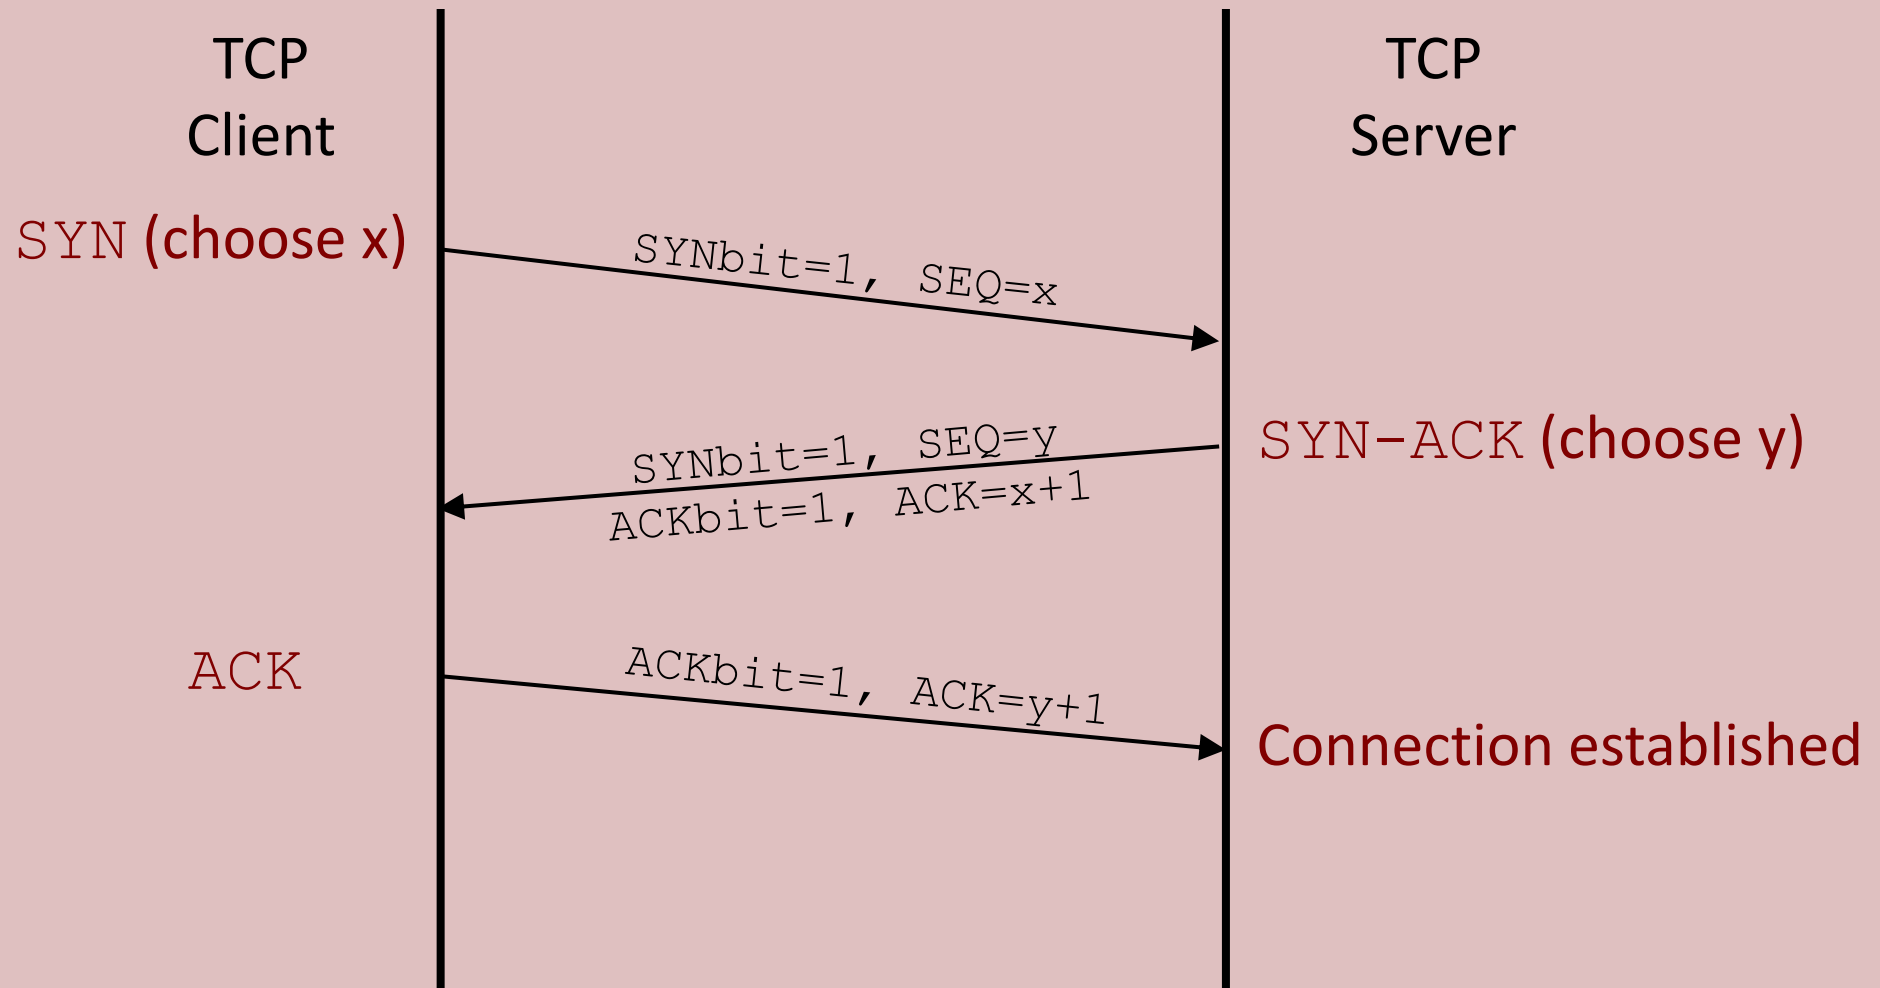
# TCP SEQ and ACK Numbers



# TCP Connections

- Perform 3-way handshake to establish the connection
  - Messages set SYN and/or ACK flags
  - Choose 32-bit initial sequence numbers
- Each side half-closes their end to terminate the connection (4 messages)
  - Messages set FIN or ACK flags

# TCP Connection Establishment



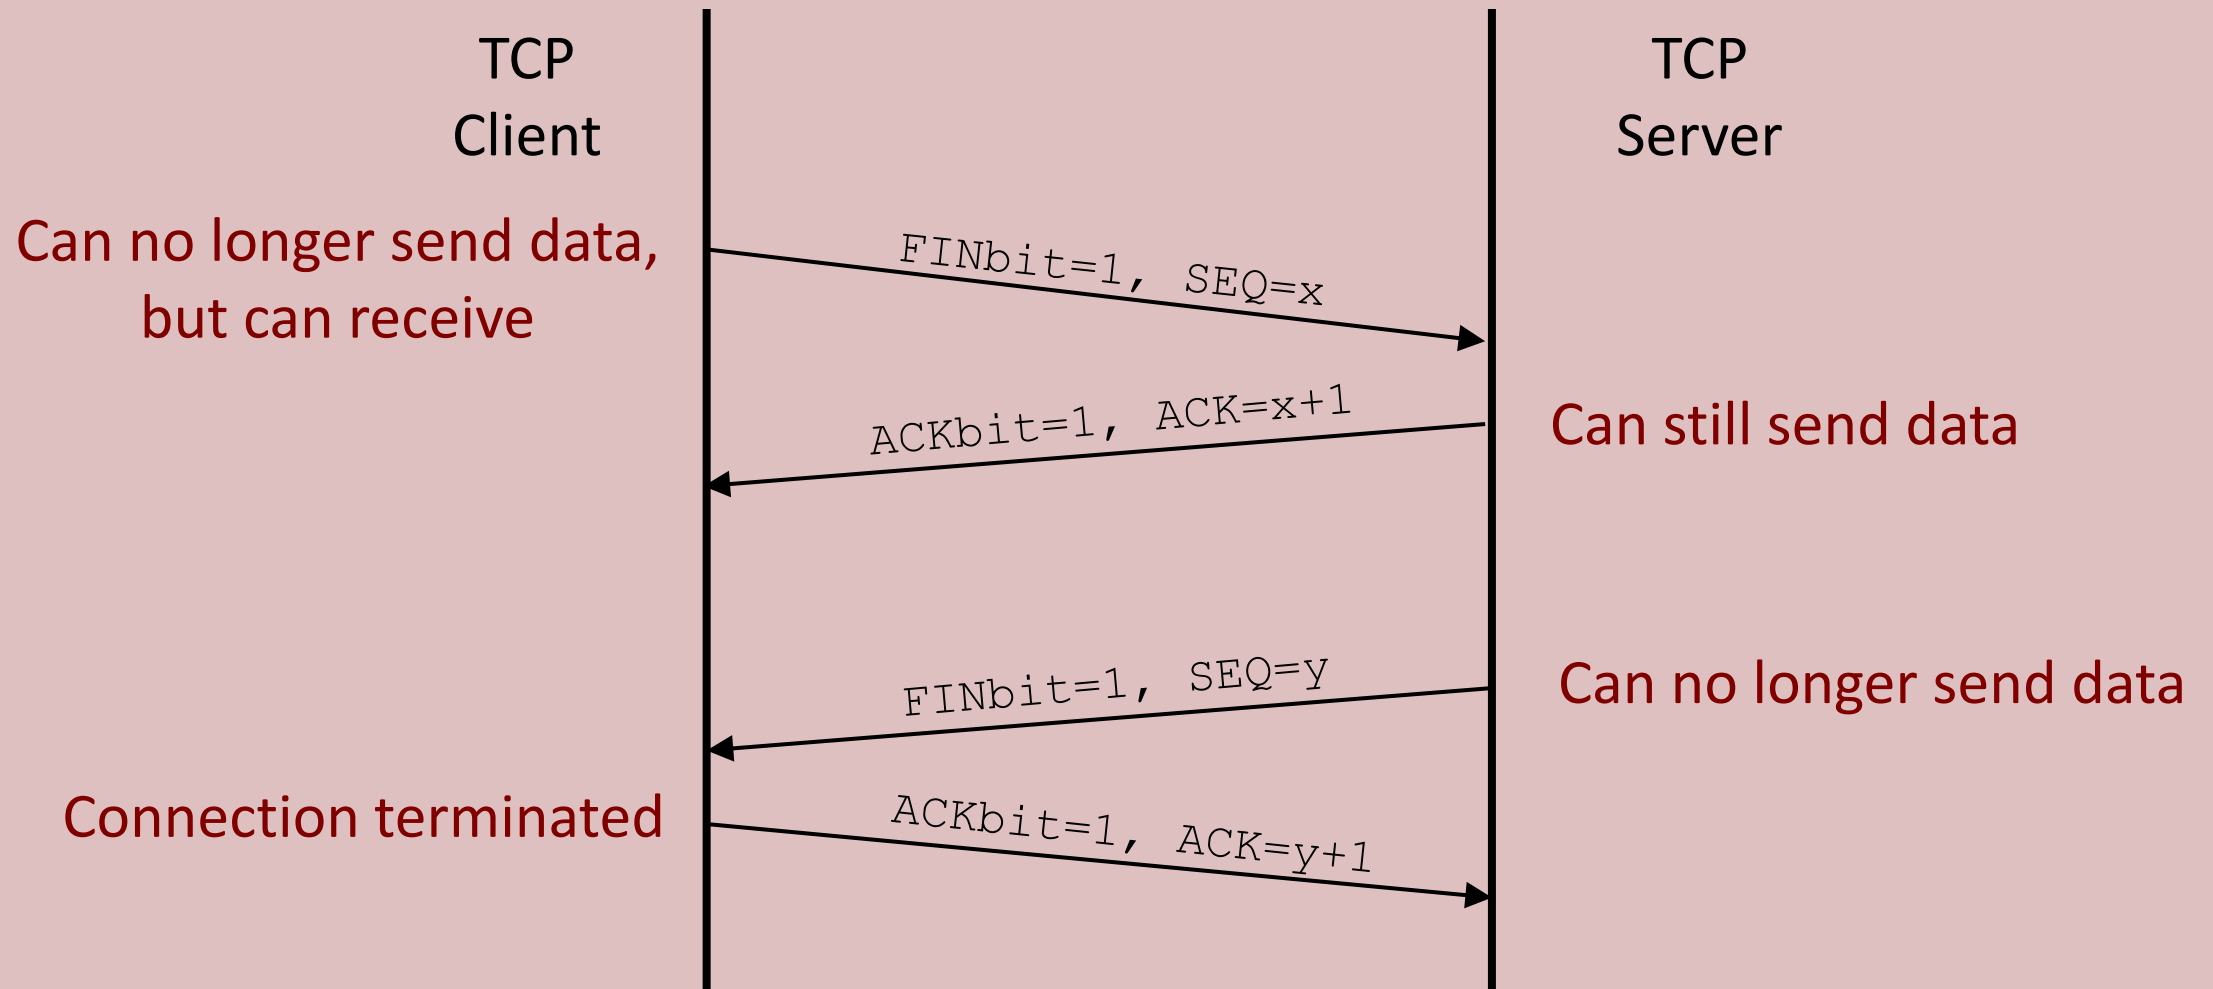
# SYN Floods

- Denial-of-service attack on TCP servers
- Attacker repeatedly sends `SYN` segments to exhaust victim server
  - Typically uses spoofed IP address
  - Server allocates resources
  - Half-open connections
- Attacker never completes 3-way handshake

# SYN Cookies

- Encode connection state in initial sequence number (ISN) chosen by server for TCP `SYN-ACK`
- Cryptographic hash of connection 4-tuple and secret value known to server
- Attackers from spoofed IP addresses cannot guess valid TCP `ACK`
  - Why not?

# TCP Connection Termination






# TCP Timeout

- If too short, unnecessary retransmissions
- If too long, slow reaction to lost segments
- Timeout interval should be longer than RTT, but RTT varies over time
- Solution: estimate RTT based on periodic RTT sample measurements
  - SampleRTT: measured time between segment transmission and ACK receipt
  - EstimatedRTT: average several recent measurements to avoid too much variation

# TCP Timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- Exponentially weighted moving average
- Influence of past samples decreases exponentially fast
- Typical value:  $\alpha=0.125$



# TCP Timeout

- Timeout interval: EstimatedRTT plus “safety margin”
  - Large variation in EstimatedRTT → larger safety margin

- Estimate SampleRTT deviation from EstimatedRTT:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

- Typical value:  $\beta=0.25$

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$


Estimated RTT

Safety Margin

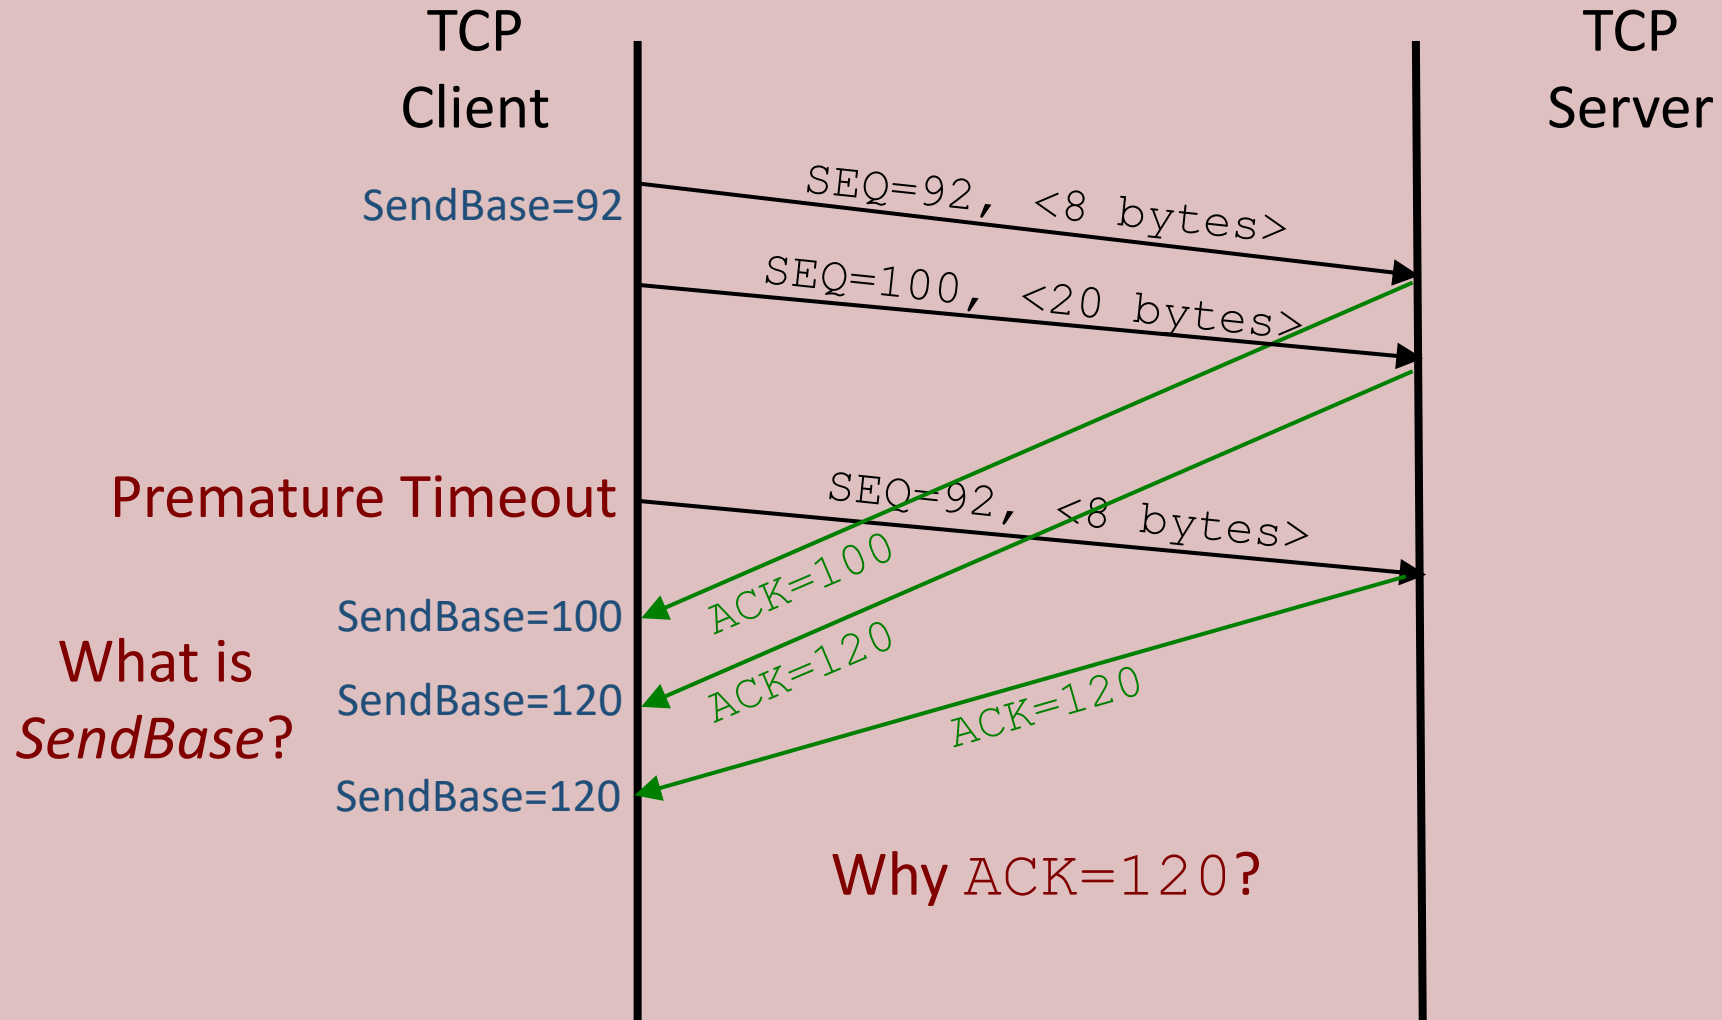
# TCP Retransmission

- Single retransmission timer that gets reset for every received ACK or timeout
- Retransmission triggered by timeout
- *Fast retransmit* also triggered by three duplicate ACKs indicating out-of-order data received

# TCP Retransmission – Lost ACK



# TCP Retransmission – Premature Timeout



# TCP Retransmission – Cumulative ACKs



# TCP ACK Generation

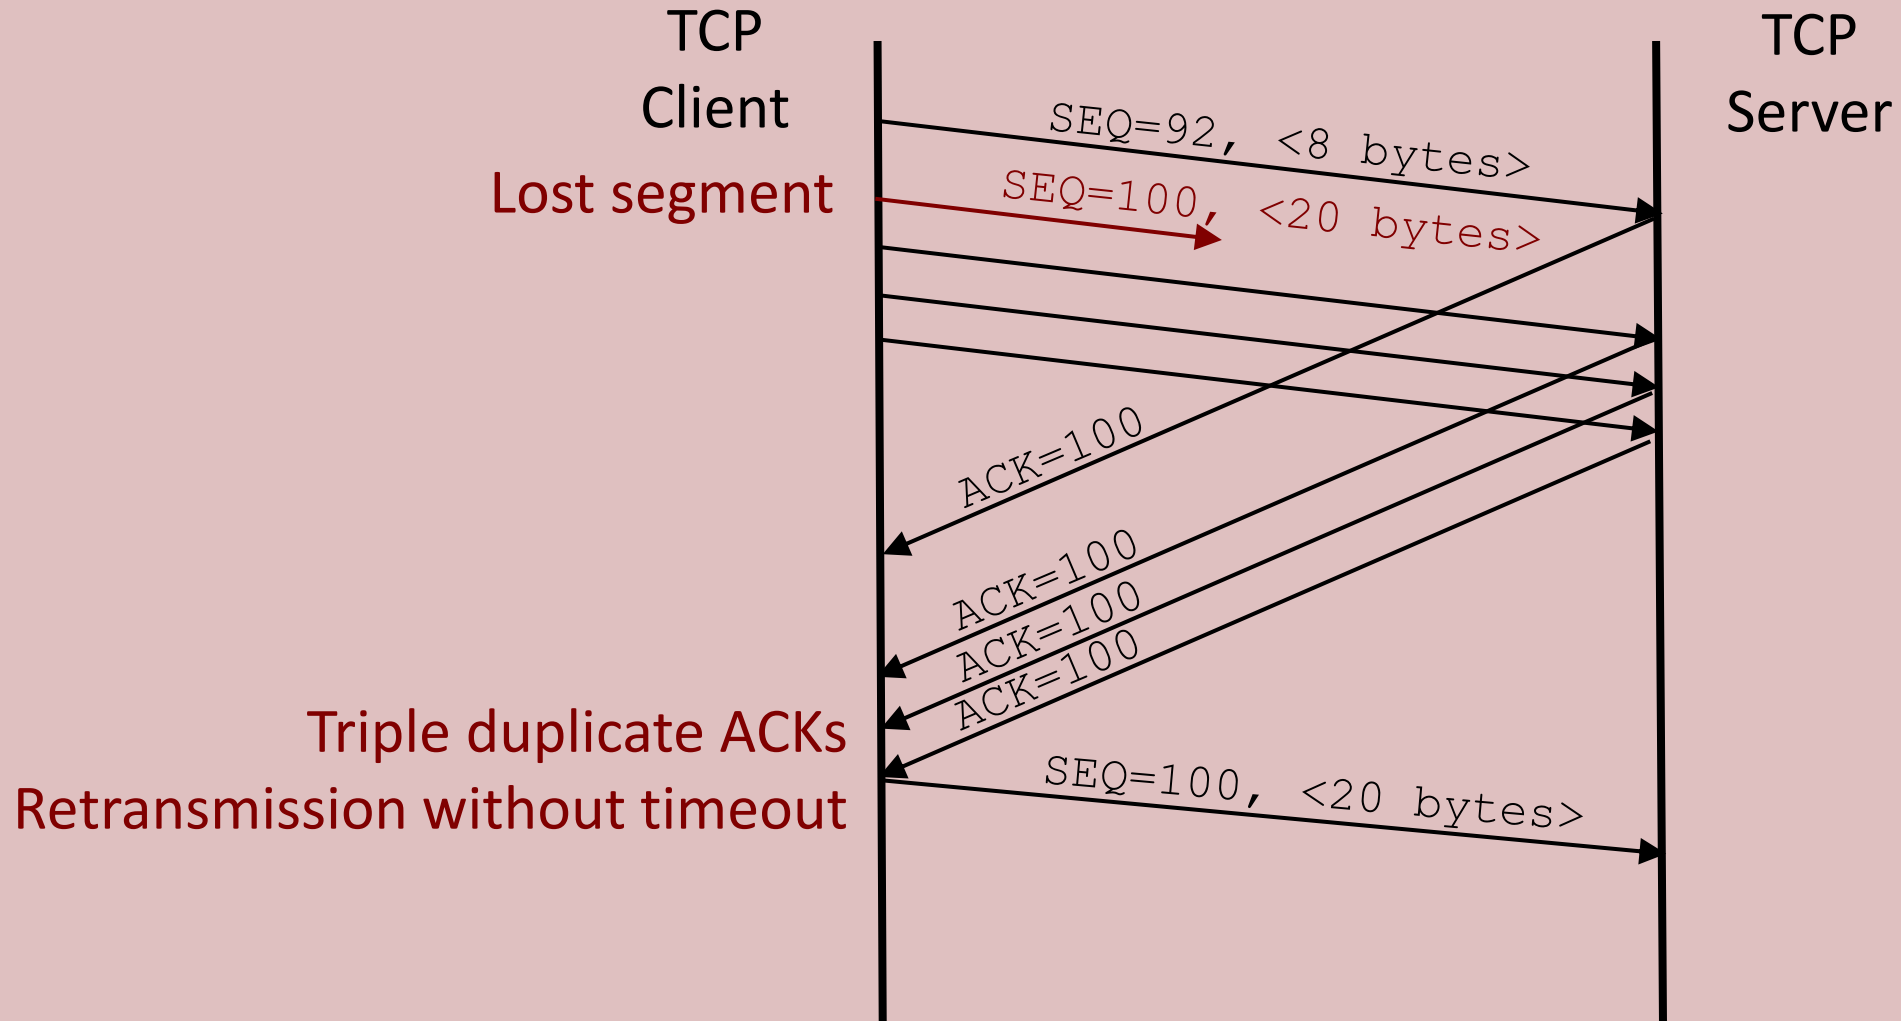
Event at Receiver	TCP Receiver Action
Arrival of in-order segment with expected SEQ #. All data up to expected SEQ # already ACKed.	Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK.
Arrival of in-order segment with expected SEQ #. One other segment has ACK pending.	Immediately send single cumulative ACK, ACKing both in-order segments.
Arrival of out-of-order segment, higher than expected SEQ #. Gap detected.	Immediately send duplicate ACK, indicating SEQ # of next expected byte.
Arrival of segment that partially or completely fills gap.	Immediately send ACK, provided that segment starts at lower end of gap.



# TCP Fast Retransmit

- Potentially long delays for timeout-based retransmissions (why?)
- Detect lost segments via duplicate ACKs
- If sender receives triple duplicate ACKs, then resend unACKed segment with smallest sequence number

# TCP Fast Retransmit



# Thank You!

# Networks

Reliable Data Transfer

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Reliable Data Transfer

## Section 3.4

# Unreliable Channels

- Packets might have bit errors
- Packets might be reordered
- Packets might be duplicated
- Packets might be dropped
- Can reliable data transfer be provided as a transport layer service?

# Automatic Repeat Request (ARQ)

- Reliable data transfer protocols that retransmit until data is finally received
- ARQ mechanisms
  - Acknowledgements
  - Timeouts
  - Sequence numbers
- ARQ types
  - Stop-and-Wait
  - Go-Back-N
  - Selective Repeat



# Acknowledgements (ACKs)

- Receiver sends signal back to sender that packet was successfully received
  - Handles lost or corrupted packets
  - Negative feedback version is a *NACK*
- Sender sends a packet, then
  - Retransmits packet if ACK not received
  - Retransmits packet if NACK received
  - Transmit next packet if ACK received

How can we determine this condition?

# Timeouts

- Unlike NACK, absence of ACK is determined implicitly
- Sender sets a timer that triggers retransmission if an ACK is not received before the timer expires
- Minimum timer expiration period should be at least one RTT

What happens if the timer fires before the RTT?

# Sequence Numbers

- Receiver needs to be able to detect duplicate packets
- Receiver needs to be able to handle out-of-order packets
- Sender assigns a unique sequence number to each packet for the receiver

# Stop-and-Wait


Sender:

- Send packet N
- Retransmit N (according to timer) until ACK received


Receiver:

- Receive packet N
- Send ACK for packet N if checksum verifies

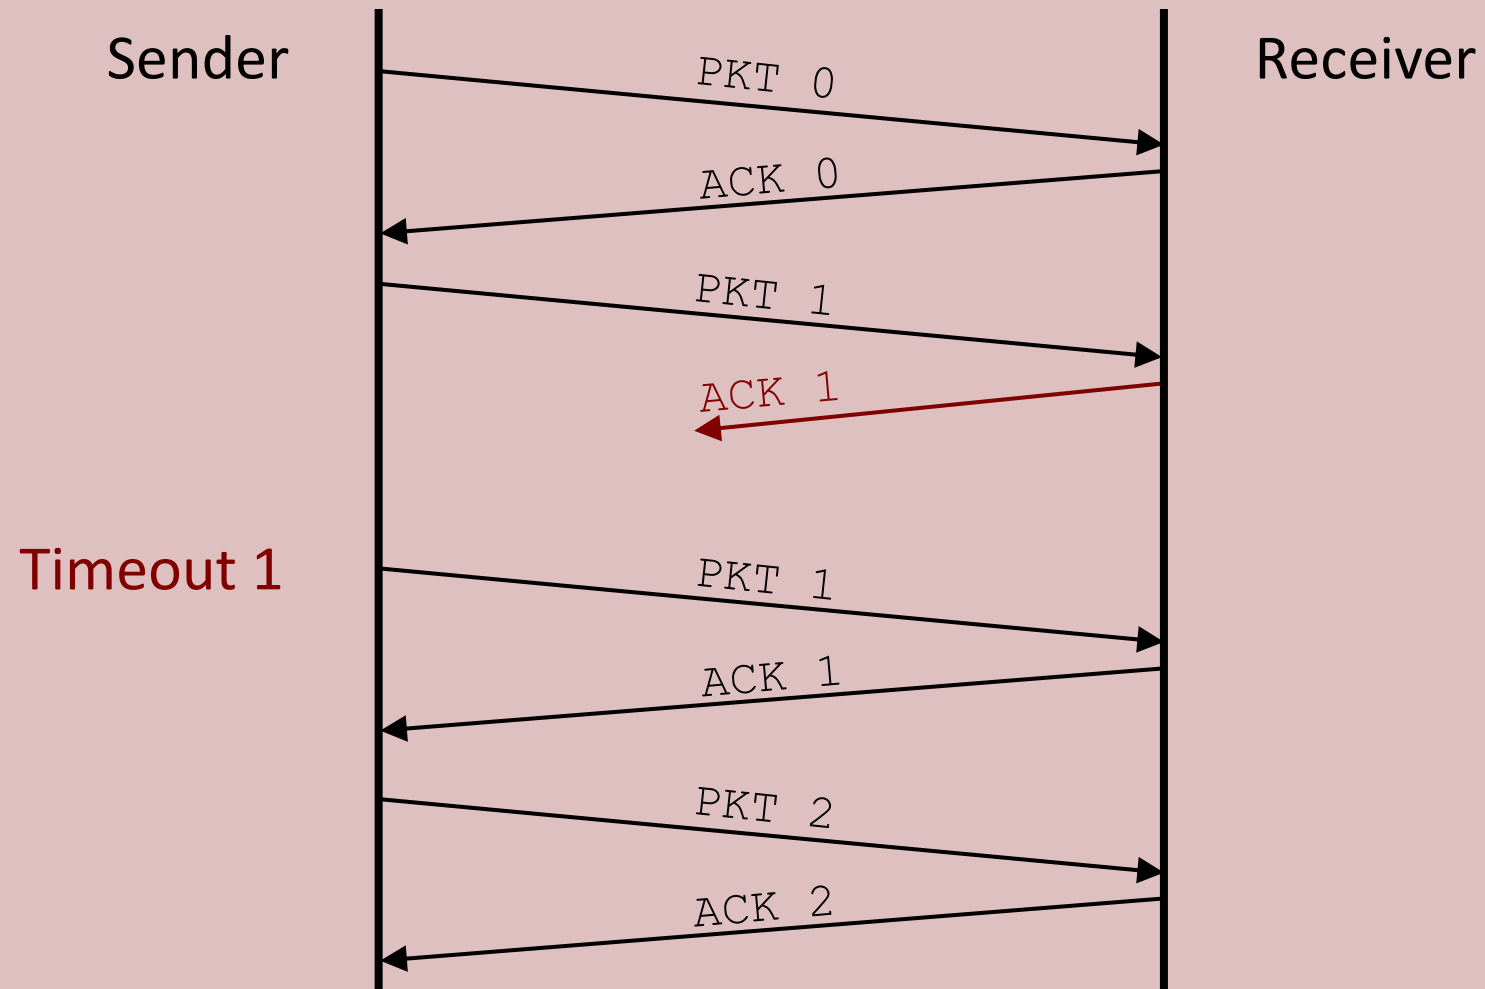
# Stop-and-Wait – No Loss



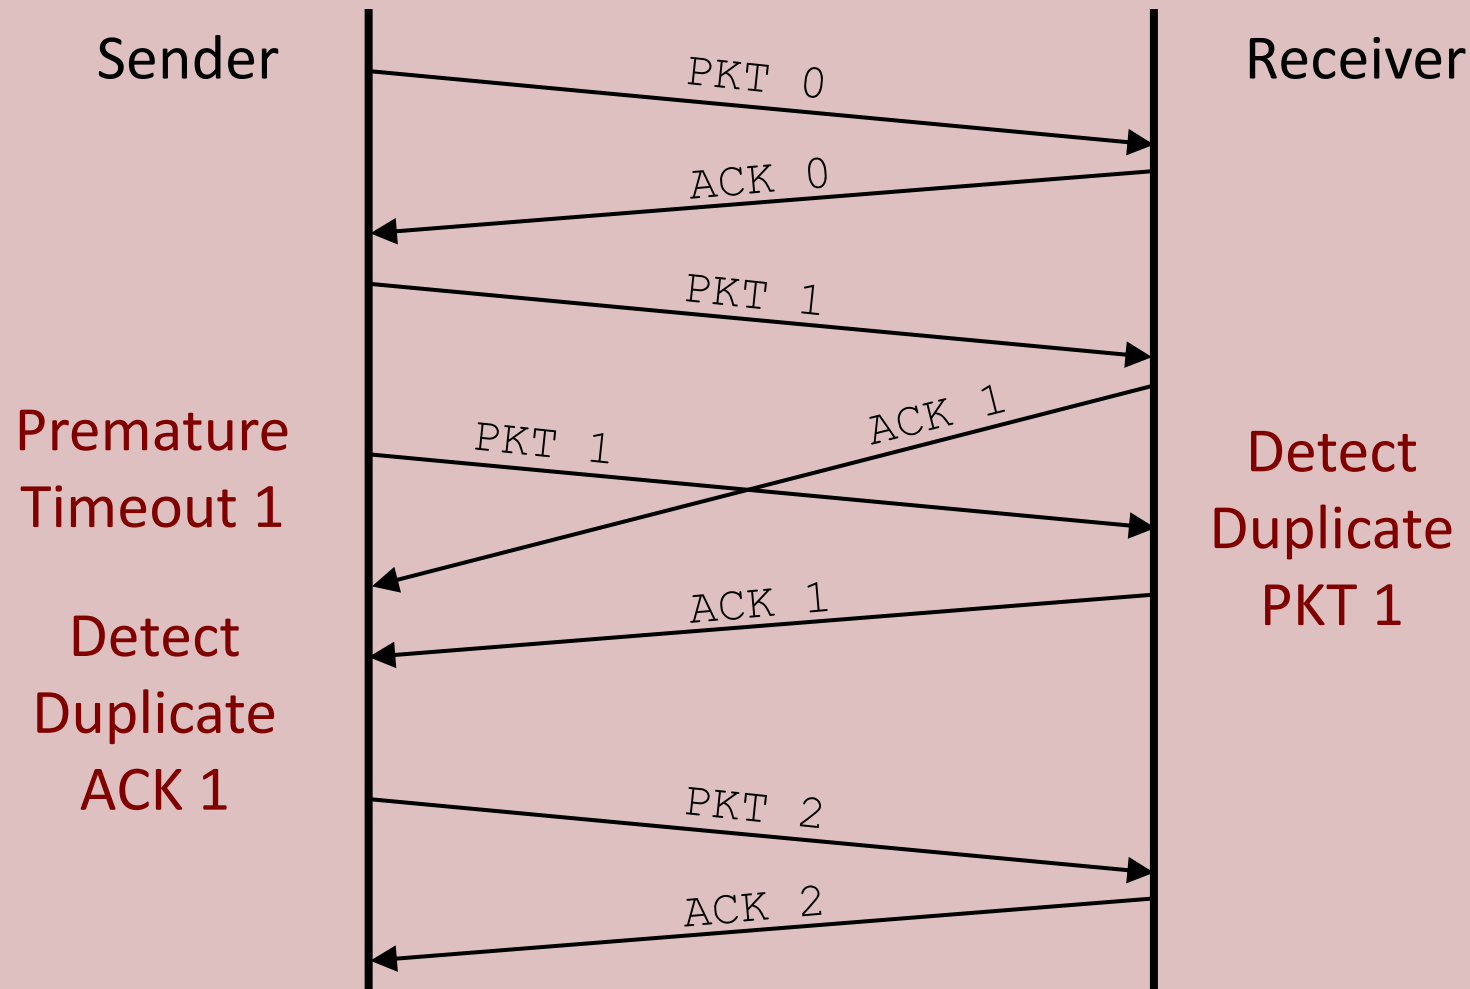
# Stop-and-Wait – Packet Loss



# Stop-and-Wait – ACK Loss



# Stop-and-Wait – Premature Timeout






# TFTP Lock Step ACKs

- Similar to two stop-and-wait protocol instances (one in each direction)
- Writer sends DATA and waits for ACKs
- Reader sends ACKs and waits for DATA (i.e., DATA is like an ACK of the ACK)
- DATA and ACKs can be retransmitted based on timers

# Stop-and-Wait



$$U_{\text{sender}} = \frac{L / R}{RTT + L / R}$$


Very poor utilization for small packets, high transmission rates, and high RTTs!

Does this problem look familiar?

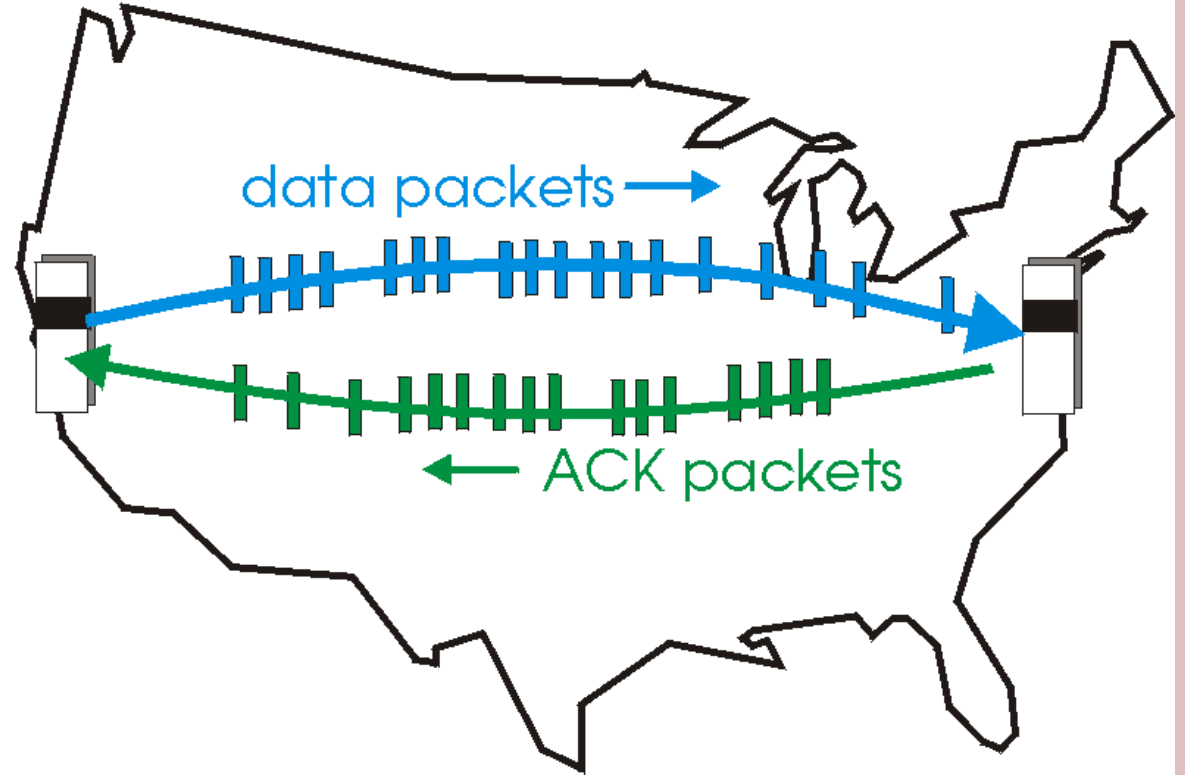
# Pipelining

- Improves efficiency over stop-and-wait
- Sender is allowed to send from a sliding window
  - Collection of packets with a subset of unacknowledged packets
  - Fixed or variable window size
- Packet buffers are required at sender and also possibly at receiver

# Pipelining

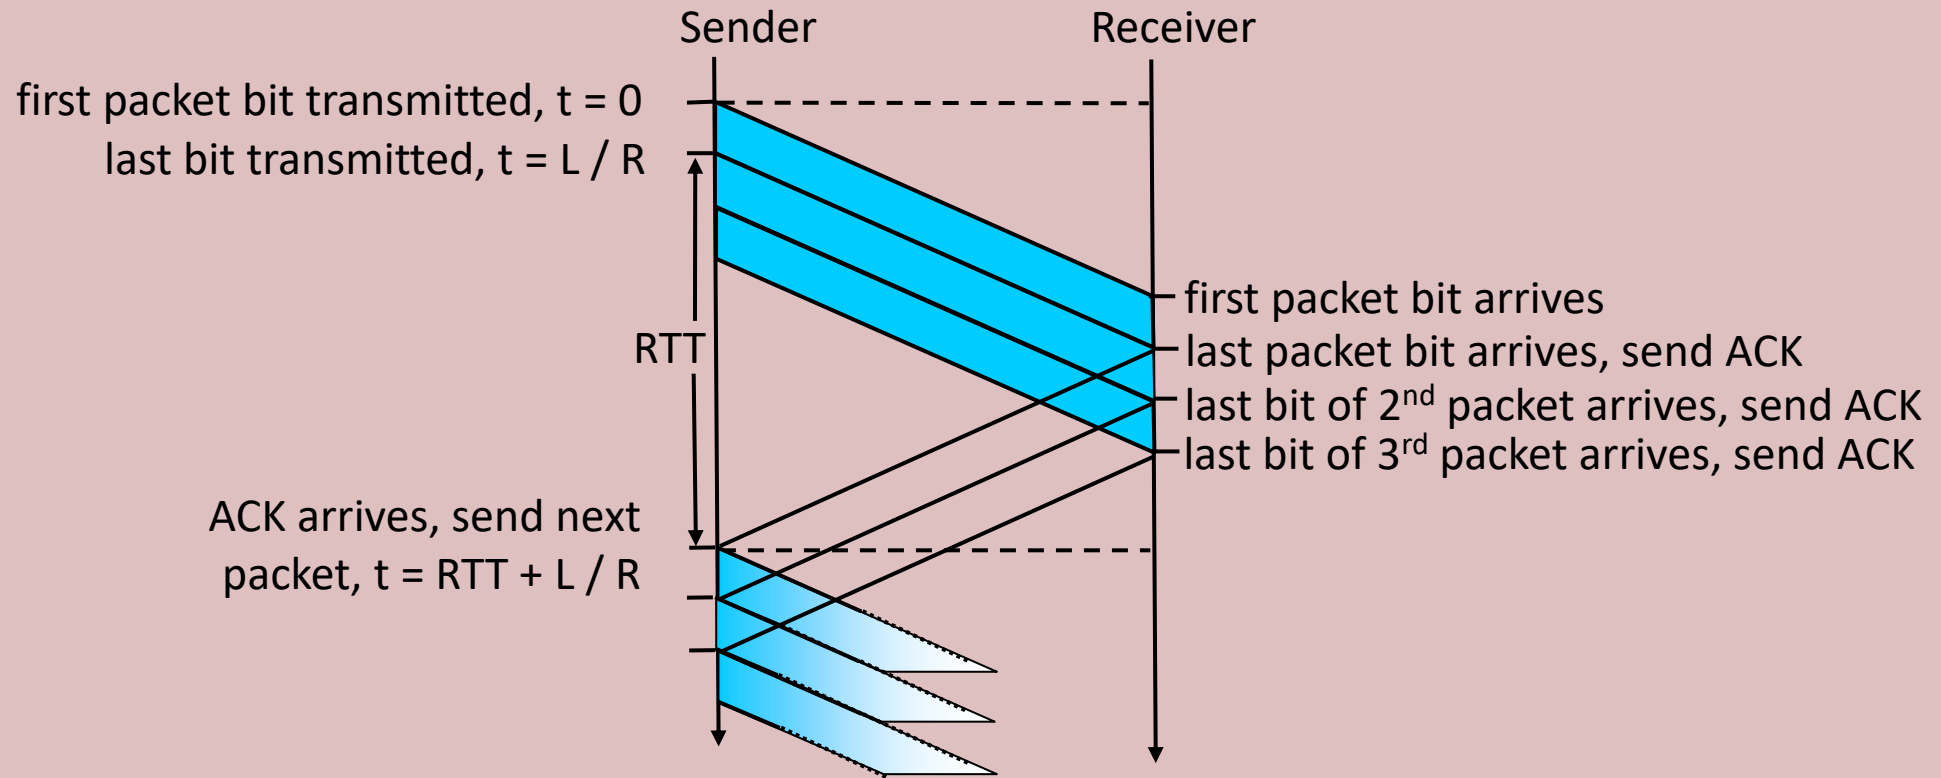


(a) a stop-and-wait protocol in operation



(b) a pipelined protocol in operation


# Pipelining




$$U_{\text{sender}} = \frac{3L / R}{RTT + L / R}$$

3-packet pipelining increases utilization by a factor of 3!


# Sliding Windows




recv(3)




recv(4)




# Sliding Windows



recvd(4)



recvd(3)





# Go-Back-N

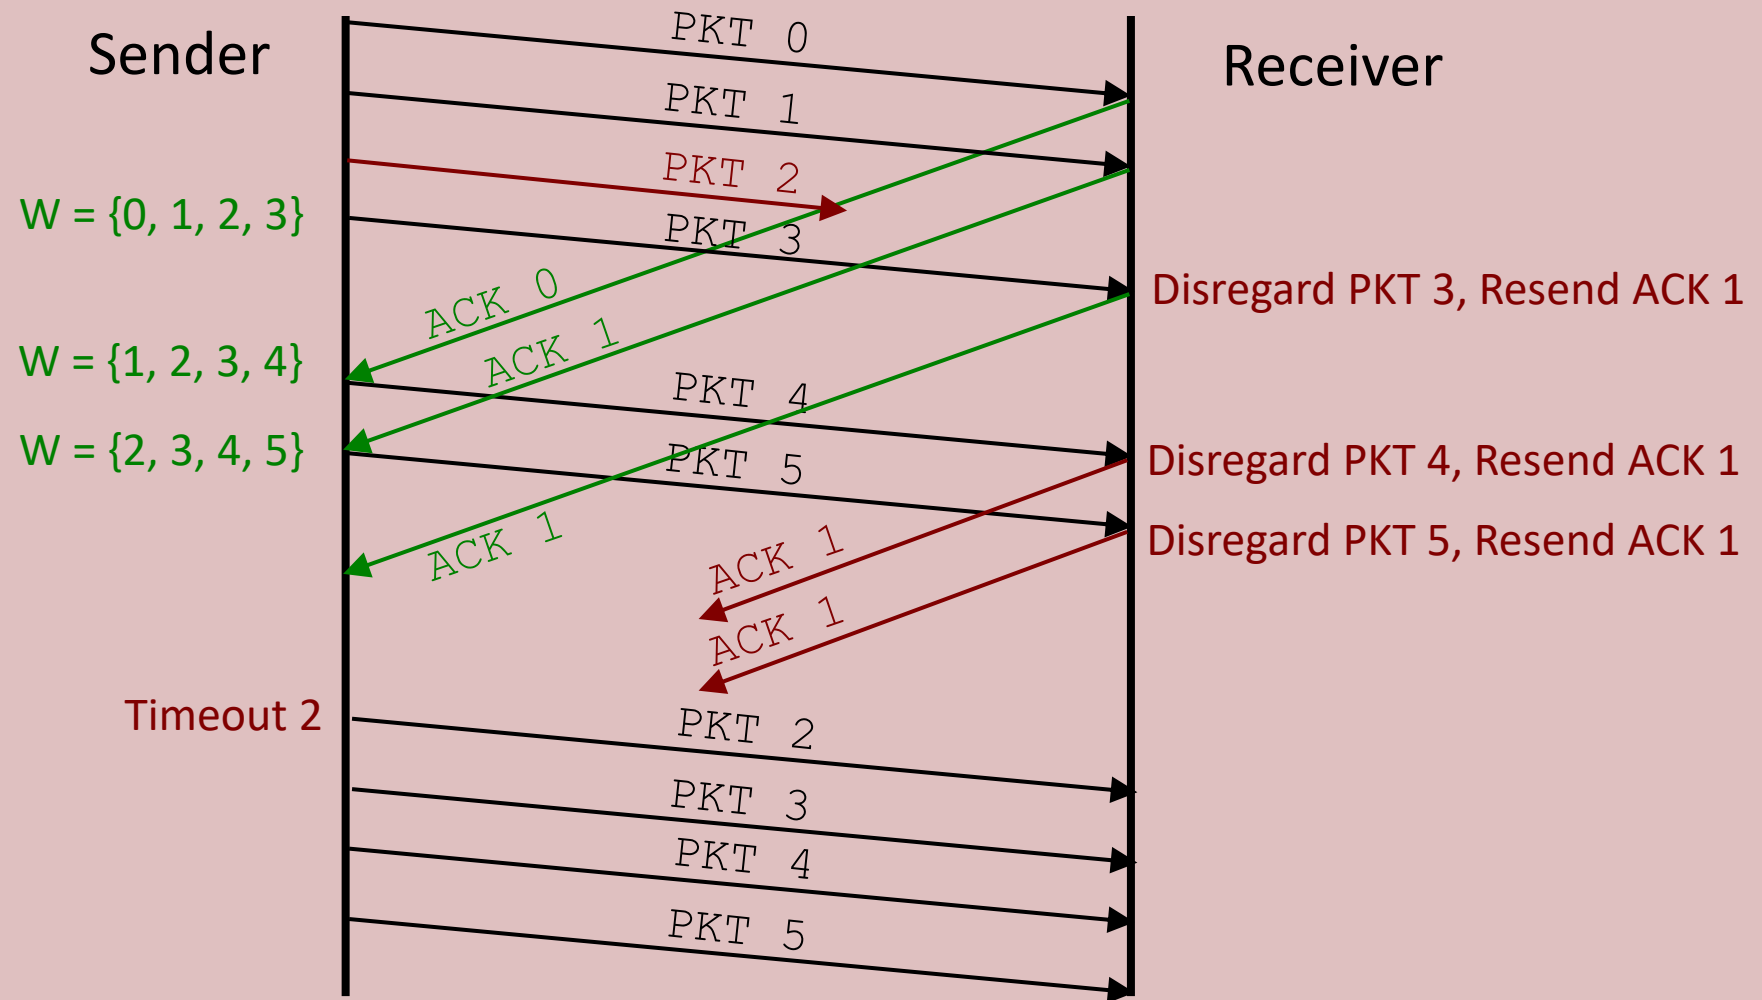
## Sender:

- Allowed to have up to  $N$  unACKed packets in pipeline
- Maintains timer for oldest unACKed packet
- Retransmit all unACKed packets when timer expires

## Receiver:

- Sends cumulative ACKs (i.e., ACK all packet numbers up to sequence number  $X$ )
- Do not ACK packet if there is a gap
- No need to buffer out-of-order packets

# Go-Back-N



# Selective Repeat


## Sender:

- Allowed to have up to  $N$  unACKed packets in pipeline
- Maintains timer for each unACKed packet
- Retransmit unACKed packet when timer expires

## Receiver:

- Sends individual ACK for each packet
- Must buffer out-of-order packets

# Selective Repeat



# Thank You!