

# Networks

## Network Layer - Introduction

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

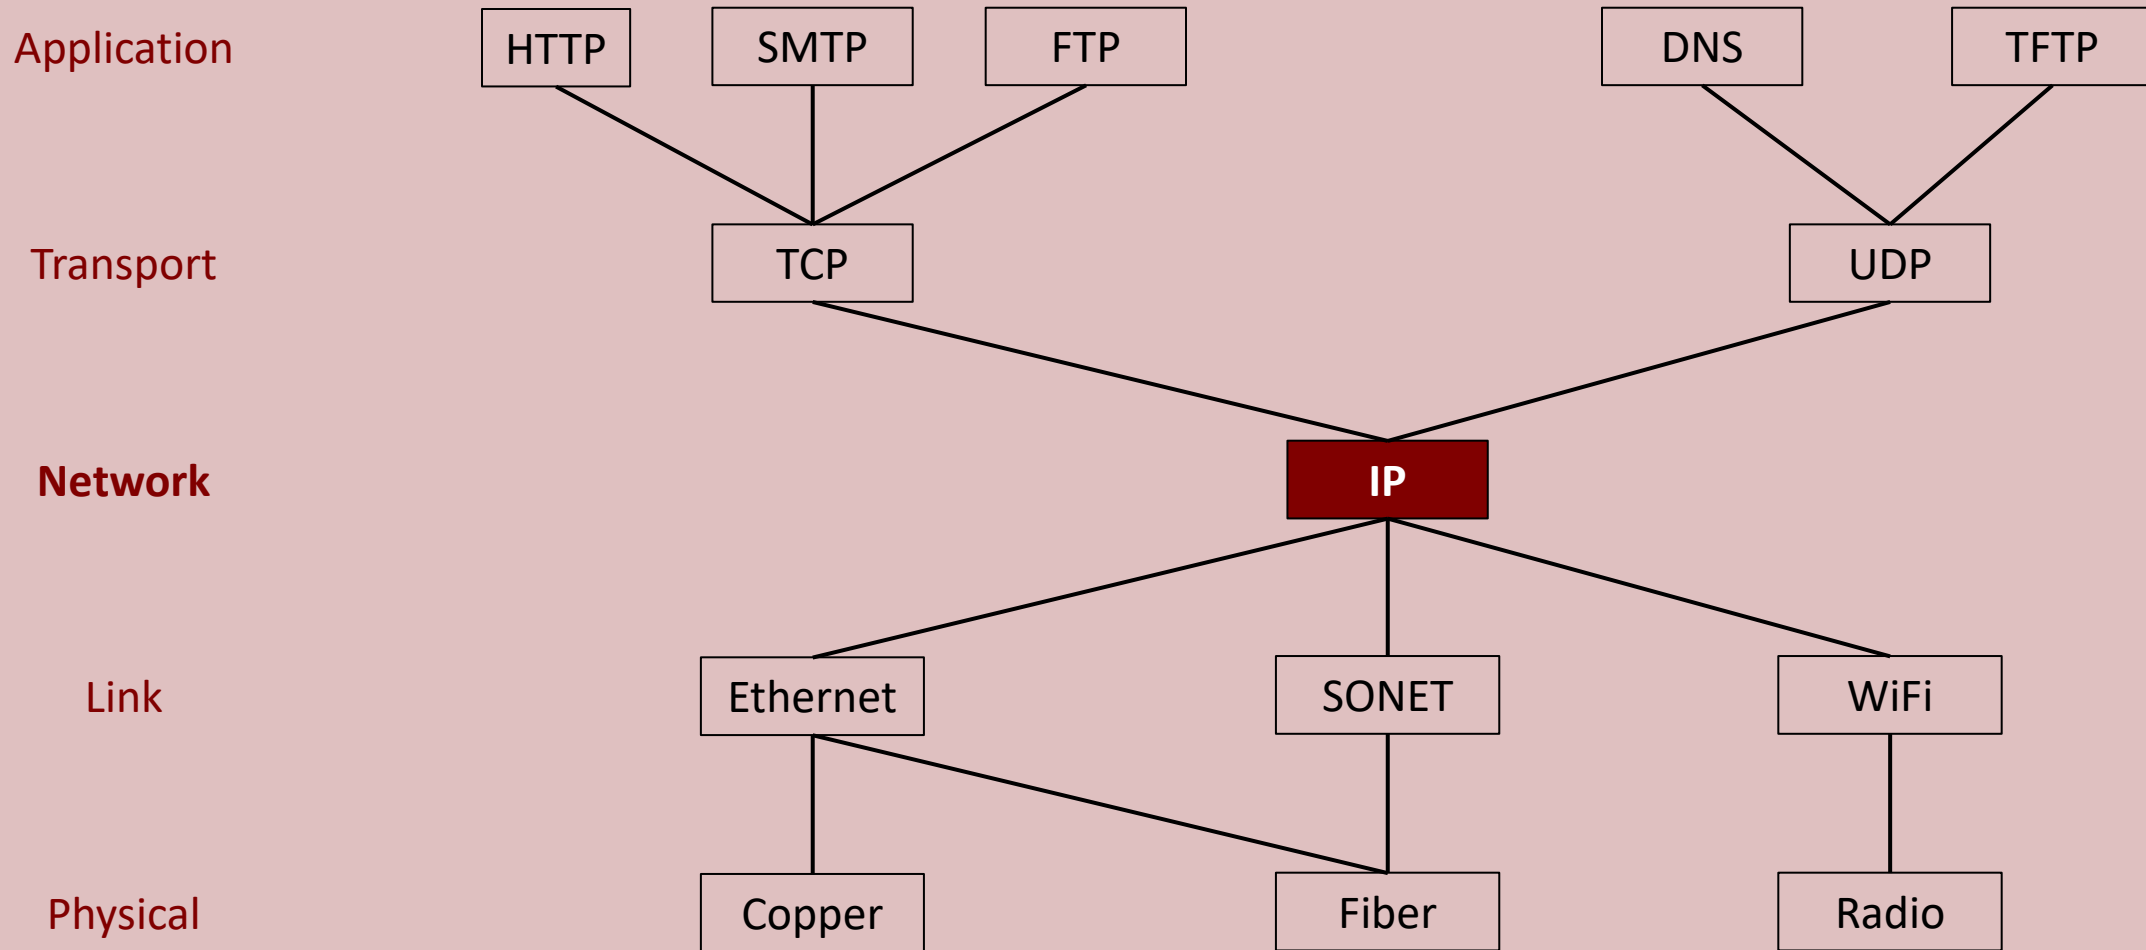
# Network Layer - Introduction

## Section 4.1

# Network Layer

- Transports either TCP segments or UDP datagrams from sending host to receiving host
  - Encapsulates transport layer data into packets at sender
  - Delivers transport layer data at receiver
- Runs on every host and router
  - IP is common layer across entire Internet
  - IP enables internetworking

# “Narrow Waist”



# Network Layer Functions

- Forwarding: move packets from router input ports to appropriate router output ports
- Routing: determine route taken by packet from source to destination

# Data vs. Control Plane

- Data plane (this week)
  - Local, per-router function
  - Forwarding
- Control plane (next week)
  - Network-wide logic
  - Routing
    - Traditional routing algorithms
    - Software-defined networking (SDN)

# Network Layer Services

Service	TCP*	IP	ATM**
Best effort		✓	✓
Reliable delivery	✓		✓
In-order delivery	✓		✓
Congestion control	✓		✓
Delay guarantees			✓
Bandwidth guarantees			✓

\* Included transport layer protocol for comparison

\*\* Actually depends on the specific service class



# Asynchronous Transfer Mode (ATM)

- Developed in the late 1980s and 1990s
- Unrelated to automated teller machines
- Telecommunications standard for carrying voice, video, and data
- Eventually lost out to IP

# ATM

- Virtual circuits over packet-switched networks
- Variations offered bandwidth, delay, loss, and ordering guarantees
- Support multiple service classes (CBR, ABR, VBR, and UBR)

# Thank You!

# Networks

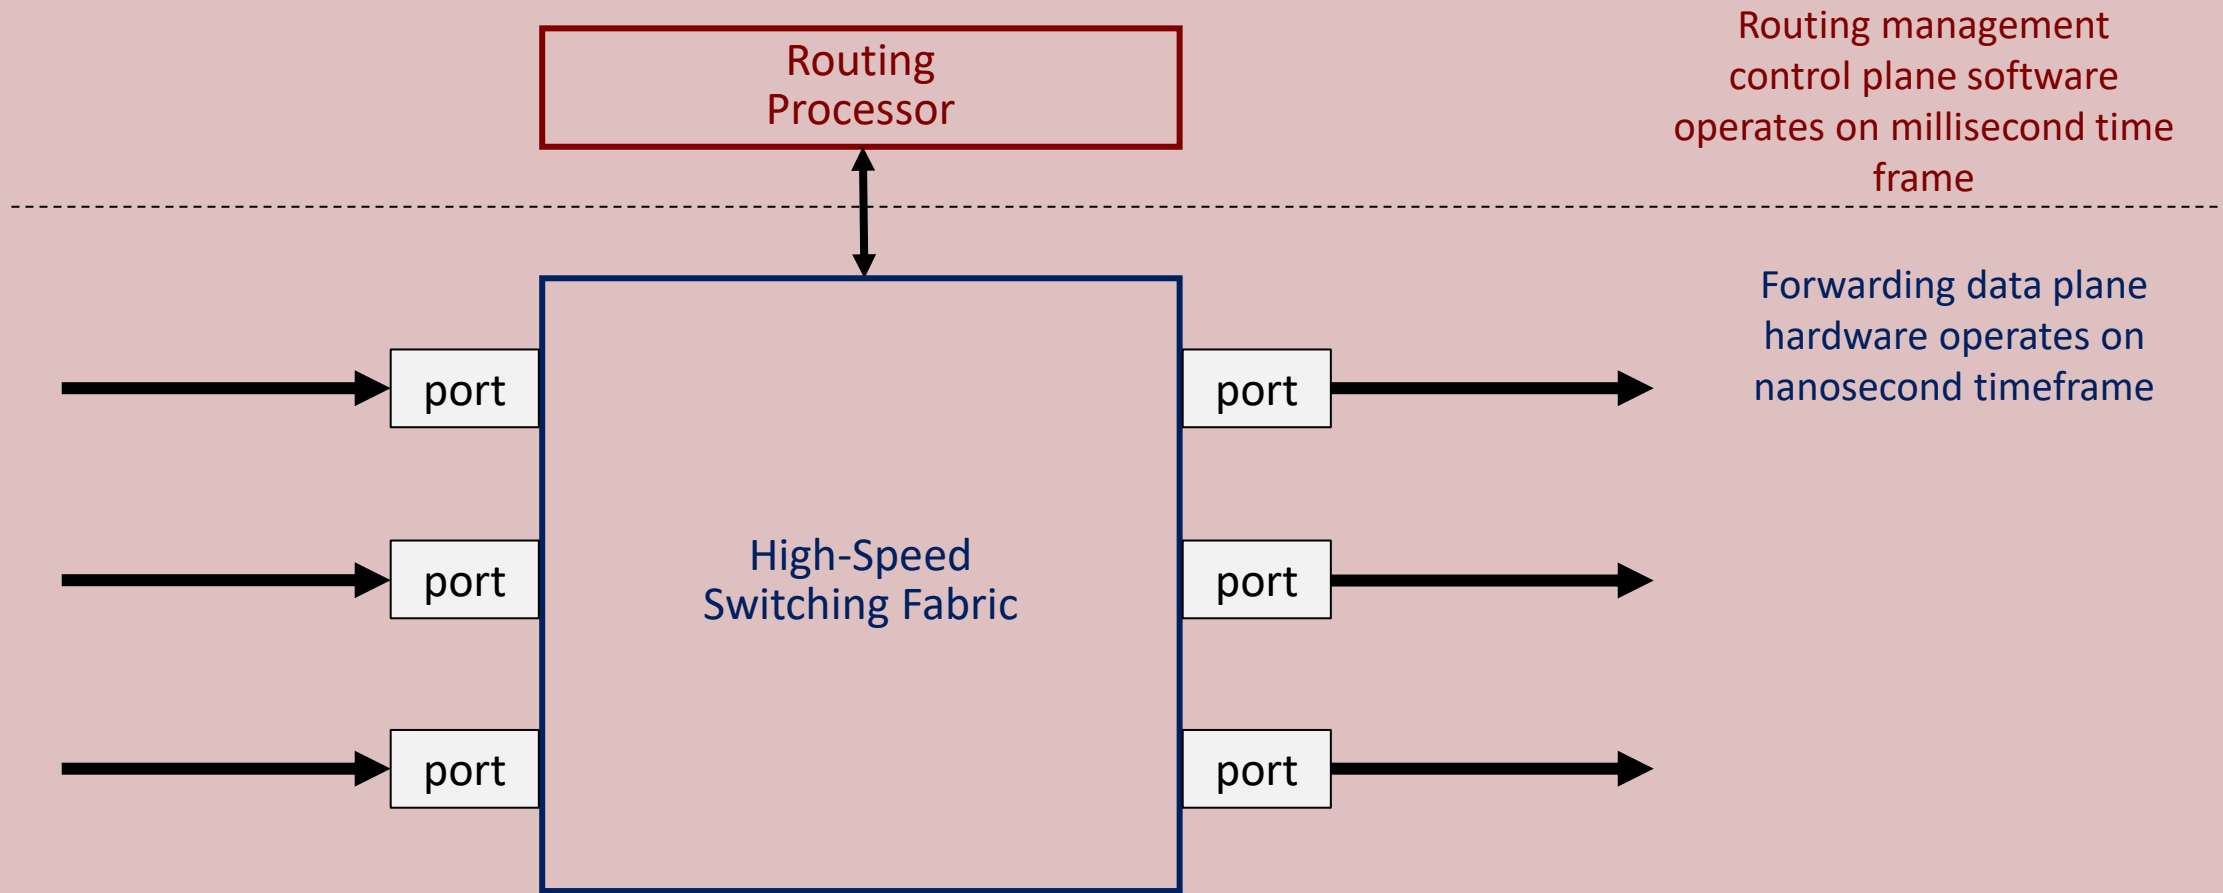
## Routers

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Routers

## Section 4.2

# Router Architecture



# Switching Functions

- Lookup
  - Determine output port using forwarding table in input port memory
- Forwarding
  - Use destination IP address (traditional)
  - Use any set of header fields (generalized)
- Queueing
  - Temporarily hold packets if arrival rate exceeds forwarding rate of switching fabric



# Destination-Based Forwarding

Forwarding Table	
Destination Address Range	Output Port
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
Default	3

What if ranges don't divide up so nicely?

# Longest Prefix Matching

Choose longest address prefix that matches destination address

Destination Address Range	Output Port
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
Default	3

Examples:

Destination Address: 11001000 00010111 00010**110** **10100001**


Destination Address: 11001000 00010111 00011**000** **10101010**

# Switching Fabric


- Transfers packets from input buffer to appropriate output buffer
- Switching rate: rate at which packets can be transferred from inputs to outputs
- Operates on the order of nanoseconds (why?)

# Switching Fabric


- Three types of switching fabric



memory




bus



crossbar


# Switching via Memory

- Used by first generation routers
- CPU directly controls switching on a traditional computer
- Packets are copied to memory
- Speed limited by memory bandwidth
  - Two bus crossings per datagram




# Switching via Bus

- Packets transferred from input port memory to output port memory via shared bus
- Switching speed limited by bus bandwidth (bus contention)
- Example: Cisco 5600
  - 32 Gbps bus
  - Sufficient speed for access and enterprise routers



# Switching via Interconnection Network

- Overcomes bus bandwidth limitations
- Originally developed to connect multiprocessors (e.g., banyan and crossbar switching)
- Advanced design
  - Fragment packet into fixed length cells
  - Switch cells through fabric
- Example: Cisco 12000
  - Switches 60 Gbps
  - “Carrier class” backbone network



# Input Port Queuing


- Can occur if switching fabric rate is lower than aggregate input line rate
- Large input queues
  - Queuing delays
  - Packet loss (input buffer overflow)

What transport layer feedback mechanism helps to reduce large input queues at routers?

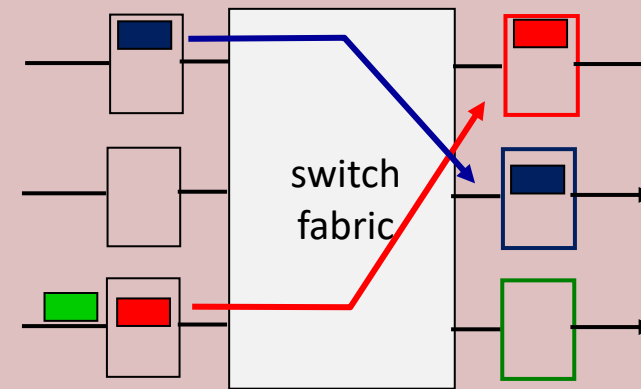


# Head-of-Line (HOL) Blocking

- Occurs if packet at front of queue prevents other queued packets from moving forward



output port contention:  
only one red datagram can be  
transferred.  
*lower red packet is blocked*

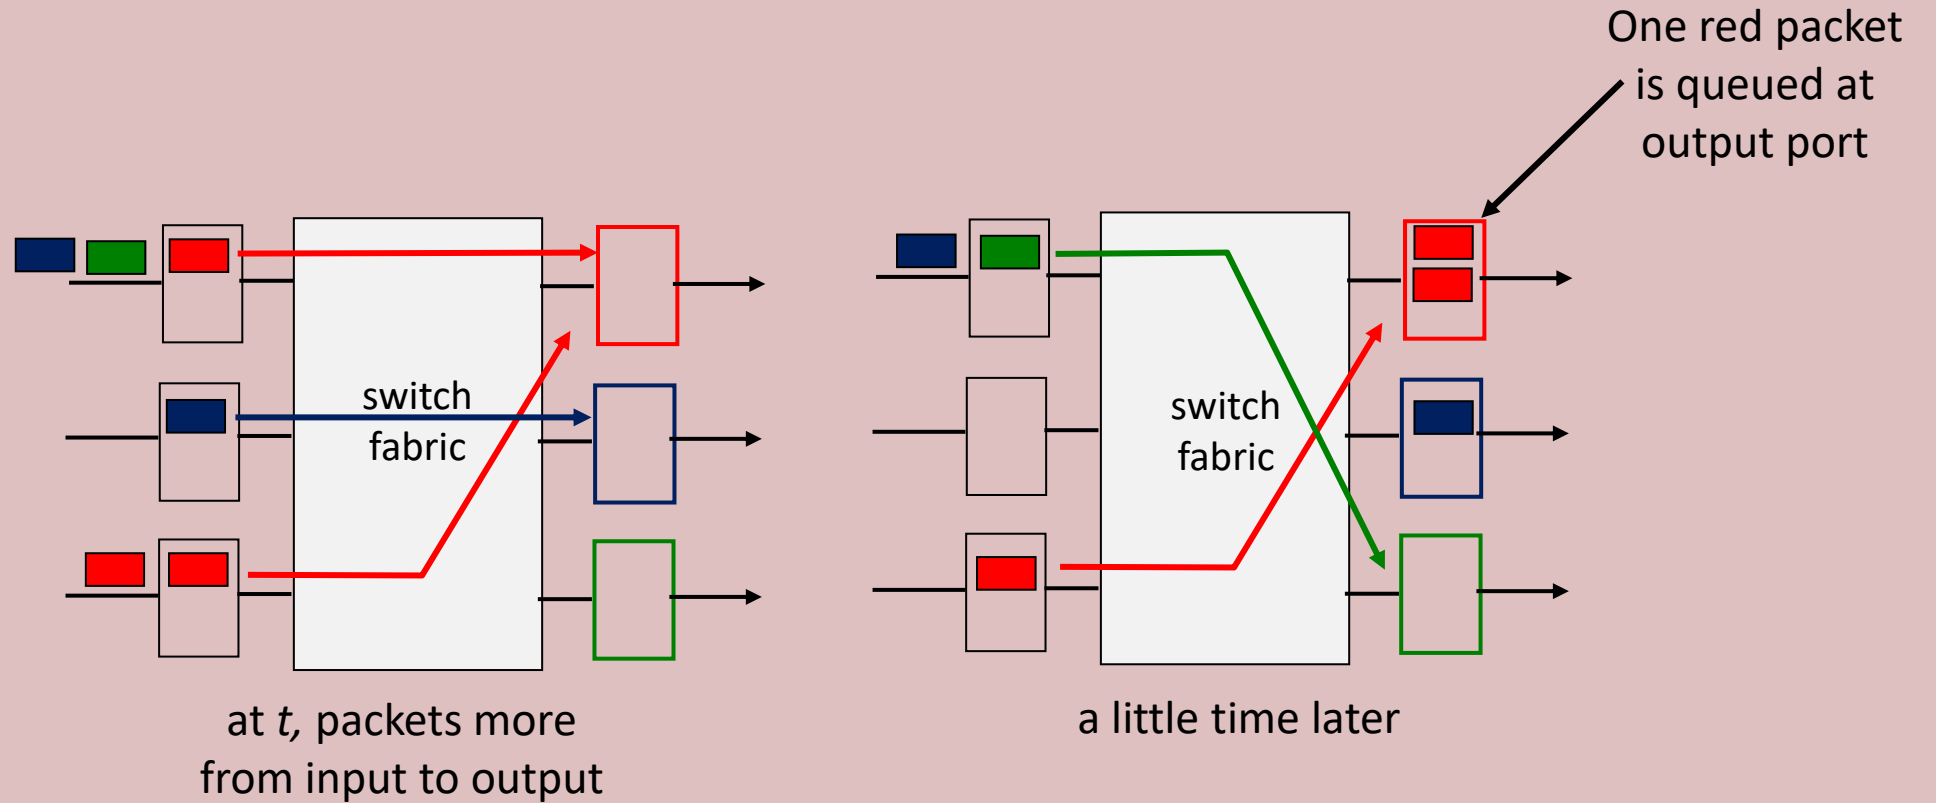


one packet time later:  
green packet  
experiences HOL  
blocking

# Output Ports


- Buffering: required when packets arrive from fabric faster than output line transmission rate
  - Queuing delays
  - Packet loss
- Scheduling discipline: choose among queued packets for next transmission on link

# Output Port Queuing



# FIFO Scheduling

- First in, first out (FIFO)
- Real-world examples
  - Sending a package at the post office
  - Boarding the bullet train in Japan
- Send packets in order of arrival
- Used in conjunction with discard policy




# Discard Policies

- Determines which packet will be dropped when new packet arrives to full queue
- Tail: drop newly arriving packet
- Priority: drop/remove on a priority basis
- Random: drop/remove random packet

# Priority Scheduling

- Send highest priority queued packet
- Packets separated by traffic classes with different priorities
  - Packet marking
  - Packet header (e.g., IP address and/or port)
- What are some real world examples of priority scheduling?



# Question

What will happen to lower priority packets during a continuous stream of higher priority packets?


# Round Robin Scheduling

- Multiple classes
- Cyclically scan queues
- Send one packet from each queue
- What are some real world examples of round robin scheduling?



# Weighted Fair Queuing

- Generalized round robin
- Each class is assigned weight proportional to amount of service received in each cycle



# Thank You!

# Networks

## Internet Protocol

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Internet Protocol

## Section 4.3

# IPv4 Datagram Header


version (4 bits)	IHL (4 bits)	diffserv (6 bits)	ECN (2)	total length (16 bits)	
identification (16 bits)				flags (3)	fragmentation offset (13 bits)
TTL (8 bits)		protocol (8 bits)		checksum (16 bits)	
source IP address (32 bits)					
destination IP address (32 bits)					
options (variable length)					

# IPv4 Datagram Header

- Version: IPv4 or IPv6
- IHL: number of 32-bit words in header
- Total length: byte length of datagram
- TTL: time to live
- Protocol: upper layer protocol
- Checksum: only protects IP header
- Fragmentation fields

# IP Fragmentation, Reassembly

- Maximum transmission unit (MTU): largest link level frame for given link type
- Large IP packet might be fragmented
- Fragments reassembled at destination
- IP header values used to identify and order the fragments





# IP Fragmentation, Reassembly

## *example:*

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

	length	ID	fragflag	offset	
	=4000	=x	=0	=0	

*one large datagram becomes  
several smaller datagrams*

1480 bytes in  
data field

offset =  $1480/8$

	length	ID	fragflag	offset	
	→ =1500	=x	=1	=0	

	length	ID	fragflag	offset	
	=1500	=x	=1	→ =185	

	length	ID	fragflag	offset	
	=1040	=x	=0	=370	

# IPv4 Addressing


- Interface: connection between host/router and physical link
  - Routers typically have multiple
  - Hosts typically have one or two (wired and/or wireless)
- IP address: 32-bit network interface identifier

# Subnet

- IP addresses consist of two parts
  - Subnet: high order bits
  - Host: low order bits
- Subnet (subnetwork)
  - Device interfaces with same subnet part of IP address
  - Can physically reach each other without an intervening router


# Subnet

Subnet Mask:  
255.255.255.0  
or /24 (shorthand)



# Subnets

- How many subnets?
- How many high order bits for subnet part?

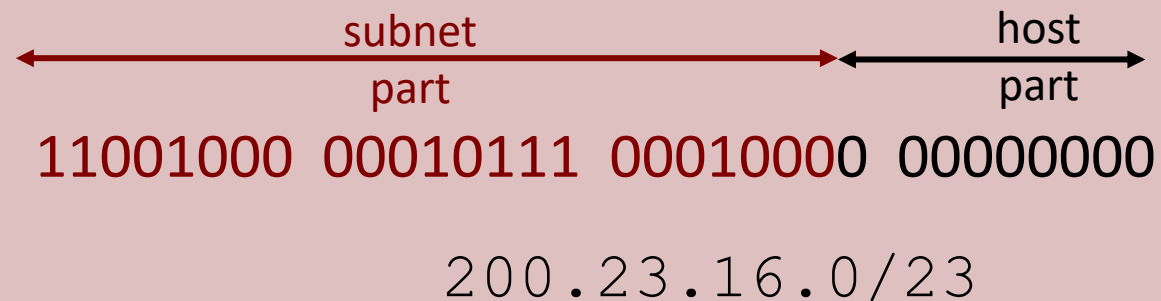


# Classful Addressing

Class	Subnet Bits	Host Bits
A	8	24
B	16	16
C	24	8
D	multicast	
E	reserved for future use	

# Classless Inter-Domain Routing (CIDR)

- Classless is more flexible and efficient than classful (why?)
- Subnet portion of address of arbitrary length
- Address format: a.b.c.d/x where x is the number of bits in the subnet portion of the address



# IP Address Assignment

- Manual configuration by sysadmin updating file
  - UNIX: /etc/rc.config
  - Windows: Control Panel > Network > Configuration > TCP/IP > Properties
- Automatic “plug-and-play” configuration via Dynamic Host Configuration Protocol (DHCP)



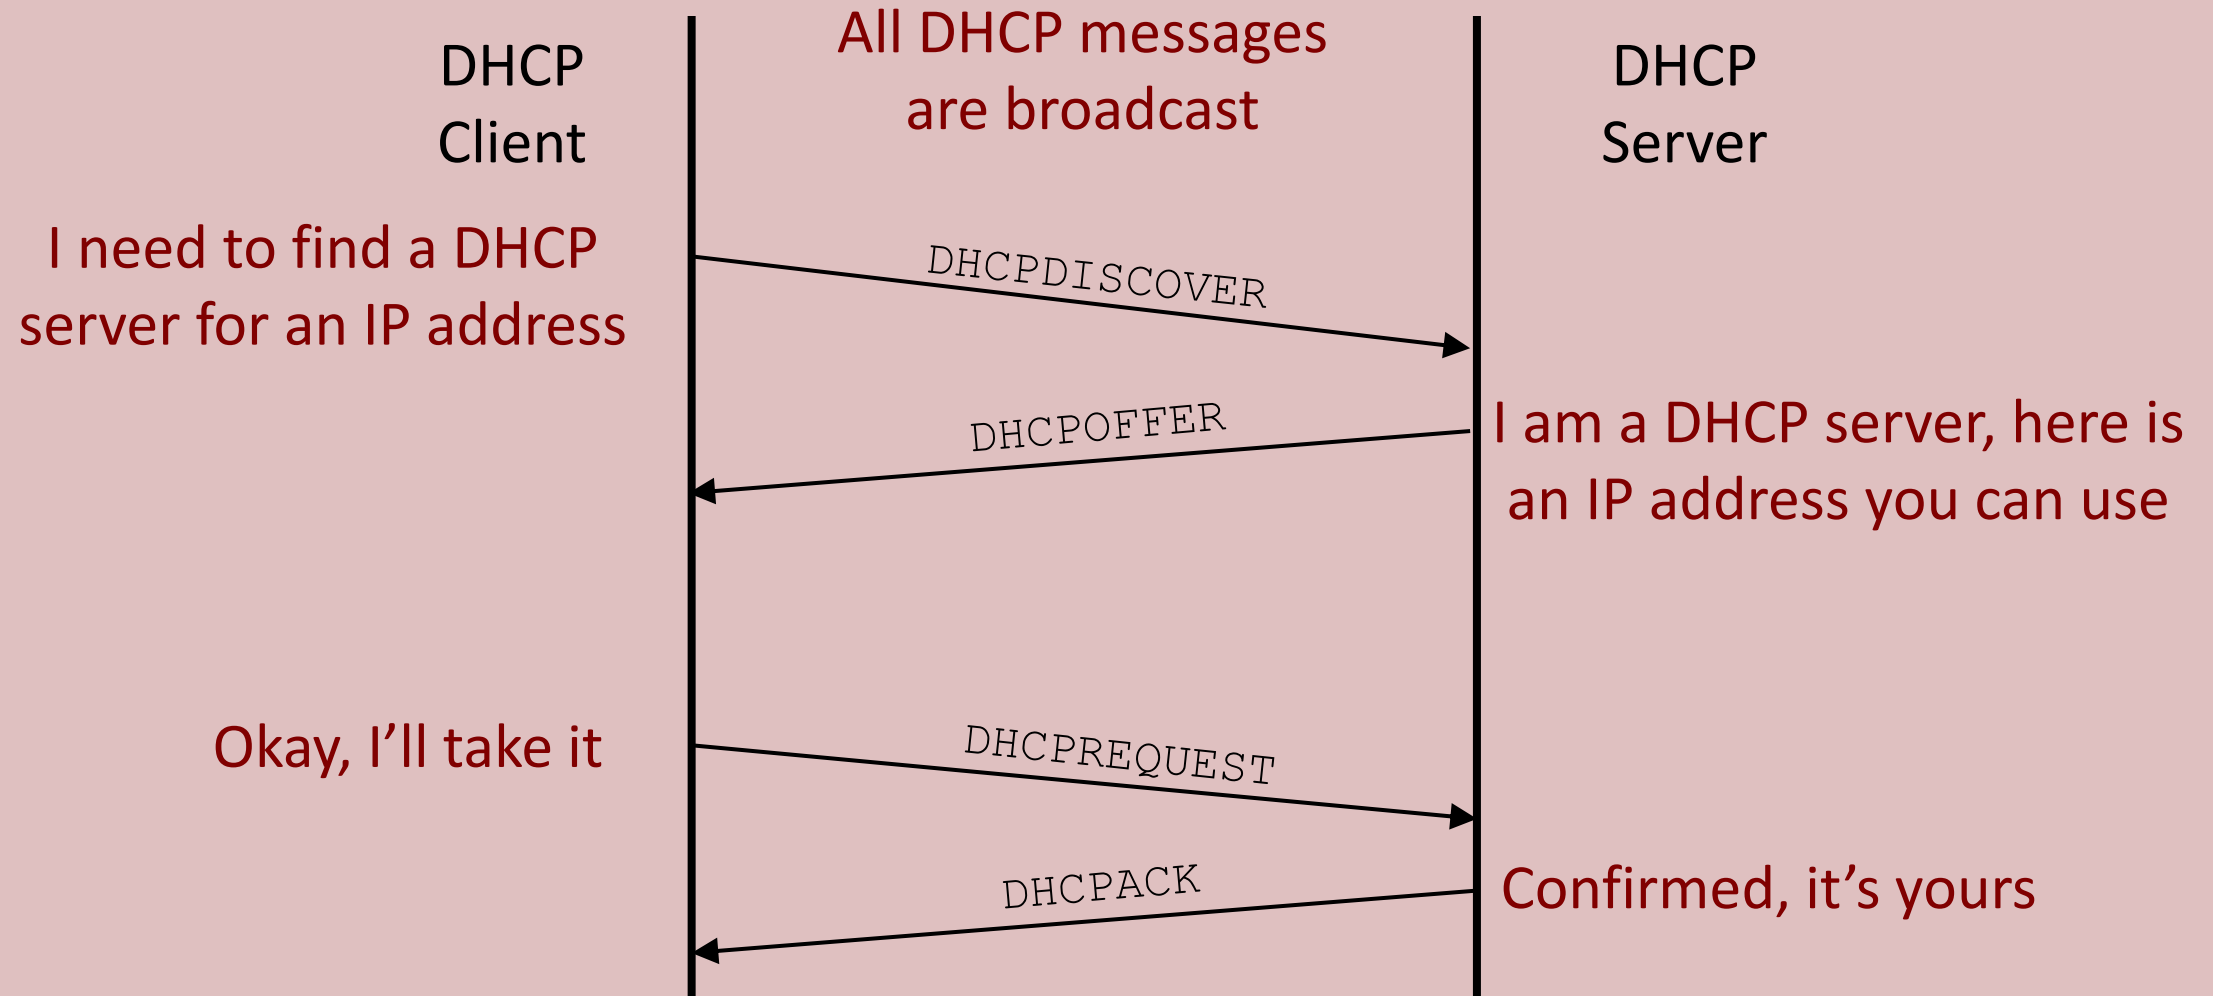
# Dynamic Host Configuration Protocol (DHCP)

- Allow host to dynamically obtain IP address from server when it joins network
- Renew lease on address in use
- Allows reuse of addresses
- Supports mobile users

# DHCP overview

- Host broadcasts DHCP discover message (optional)
- DHCP server responds with DHCP offer message (optional)
- Host requests IP address via DHCP request message
- DHCP server sends assigned IP address via DHCP ack message

# DHCP



# DHCPDISCOVER

// no client IP yet, clients use port 68

src: 0.0.0.0, 68

// broadcast message, servers use port 67

dst: 255.255.255.255, 67

// “your” IP address, not yet filled in by server

yiaddr: 0.0.0.0

// 32-bit transaction ID for exchange

xid: 654

# DHCP OFFER

// server IP and port

src: 223.1.2.4, 67

// broadcast message, clients use port 68

dst: 255.255.255.255, 68

// IP address offered by server

yiaddr: 223.1.2.5

// 32-bit transaction ID for request-response

xid: 654

// lease lifetime (in seconds)

lifetime: 3600

# DHCPREQUEST

// still no client IP yet, clients use port 68

src: 0.0.0.0, 68

// broadcast message, servers use port 67

dst: 255.255.255.255, 67

// IP address offered by server

yiaddr: 223.1.2.5

// 32-bit transaction ID for request-response

xid: 655

// lease lifetime (in seconds)

lifetime: 3600

# DHCPACK

// server IP and port

src: 223.1.2.4, 67

// broadcast message, clients use port 68

dst: 255.255.255.255, 68

// IP address offered by server

yiaddr: 223.1.2.5

// 32-bit transaction ID for request-response

xid: 655

// lease lifetime (in seconds)

lifetime: 3600

# Additional Configuration via DHCP

- DHCP can return more than just allocated IP address on subnet
- Address of first-hop router for client
- Name and IP address of DNS server
- Network mask (specify network vs. host portion of IP address)



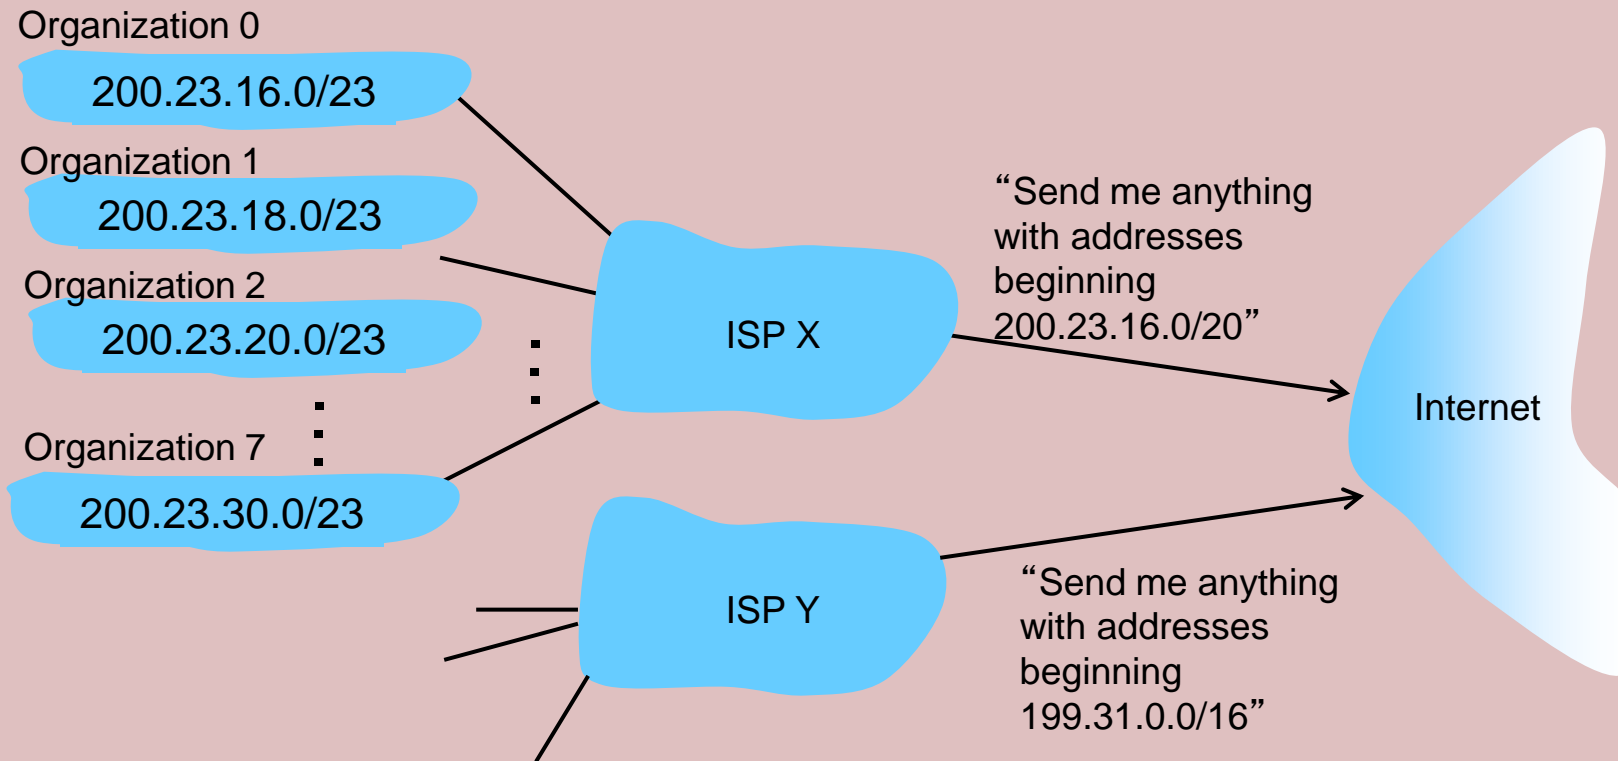
# Obtaining IP Address Blocks

- ISP allocates a portion of its own IP address space to customer organizations

ISP's block	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/20
Organization 0	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/23
Organization 1	<u>11001000</u>	<u>00010111</u>	<u>00010010</u>	00000000	200.23.18.0/23
Organization 2	<u>11001000</u>	<u>00010111</u>	<u>00010100</u>	00000000	200.23.20.0/23
...	.....				....
Organization 7	<u>11001000</u>	<u>00010111</u>	<u>00011110</u>	00000000	200.23.30.0/23

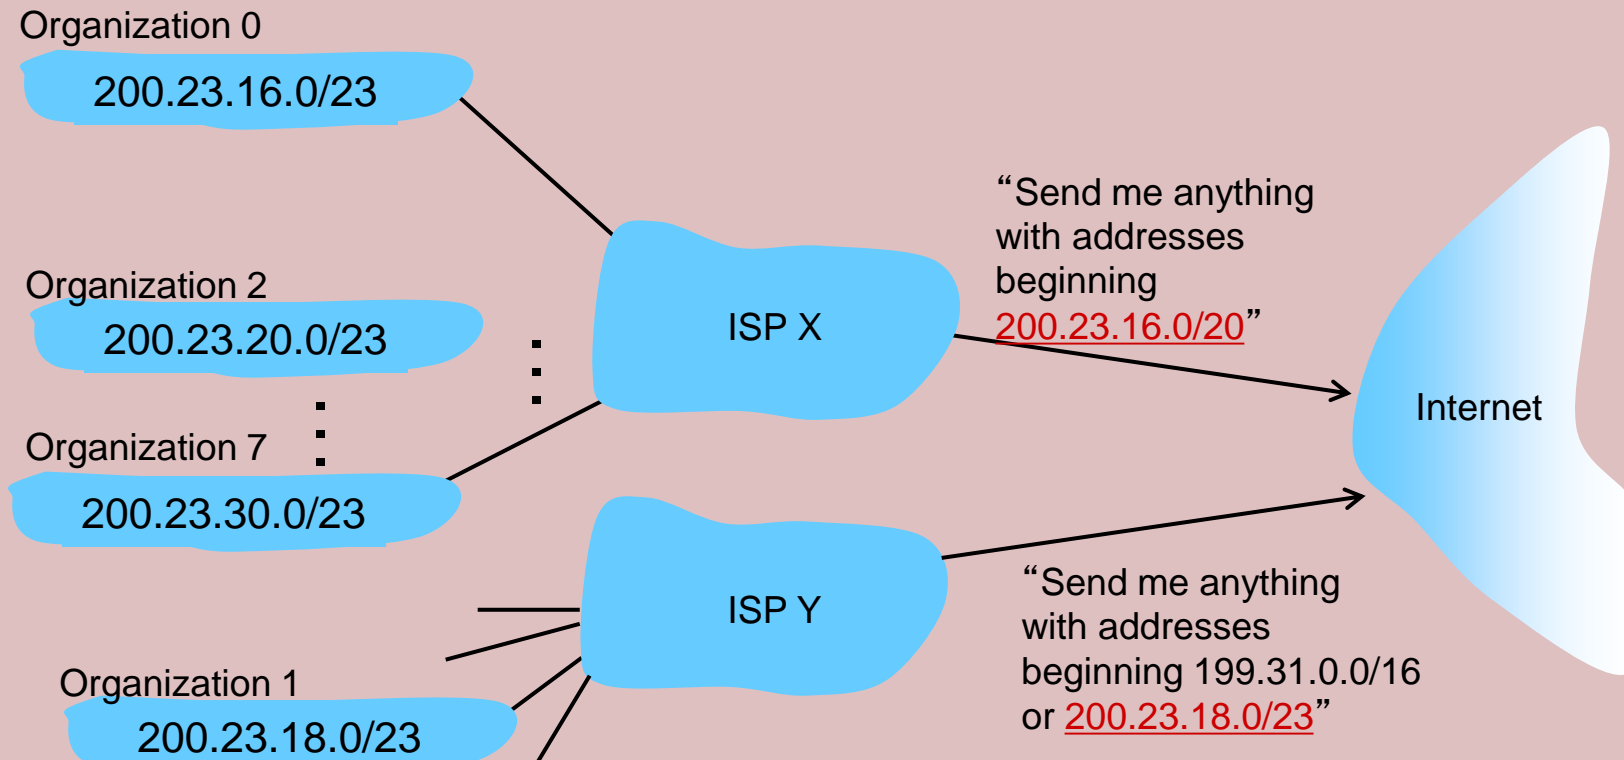
# Route Aggregation

Hierarchical addressing allows efficient advertisement of routing information:



# More Specific Routes

ISP Y has a more specific route to Organization 1




Why doesn't range overlap cause conflict?

# ICANN

- Internet Corporation for Assigned Names and Numbers (<https://www.icann.org>)
- Allocates blocks of IP addresses to ISPs
- Manages DNS
- Assigns domain names and resolves disputes

# Network Address Translation (NAT)



*all* datagrams *leaving* local network have *same* single source NAT IP address: 138.76.29.7, different source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

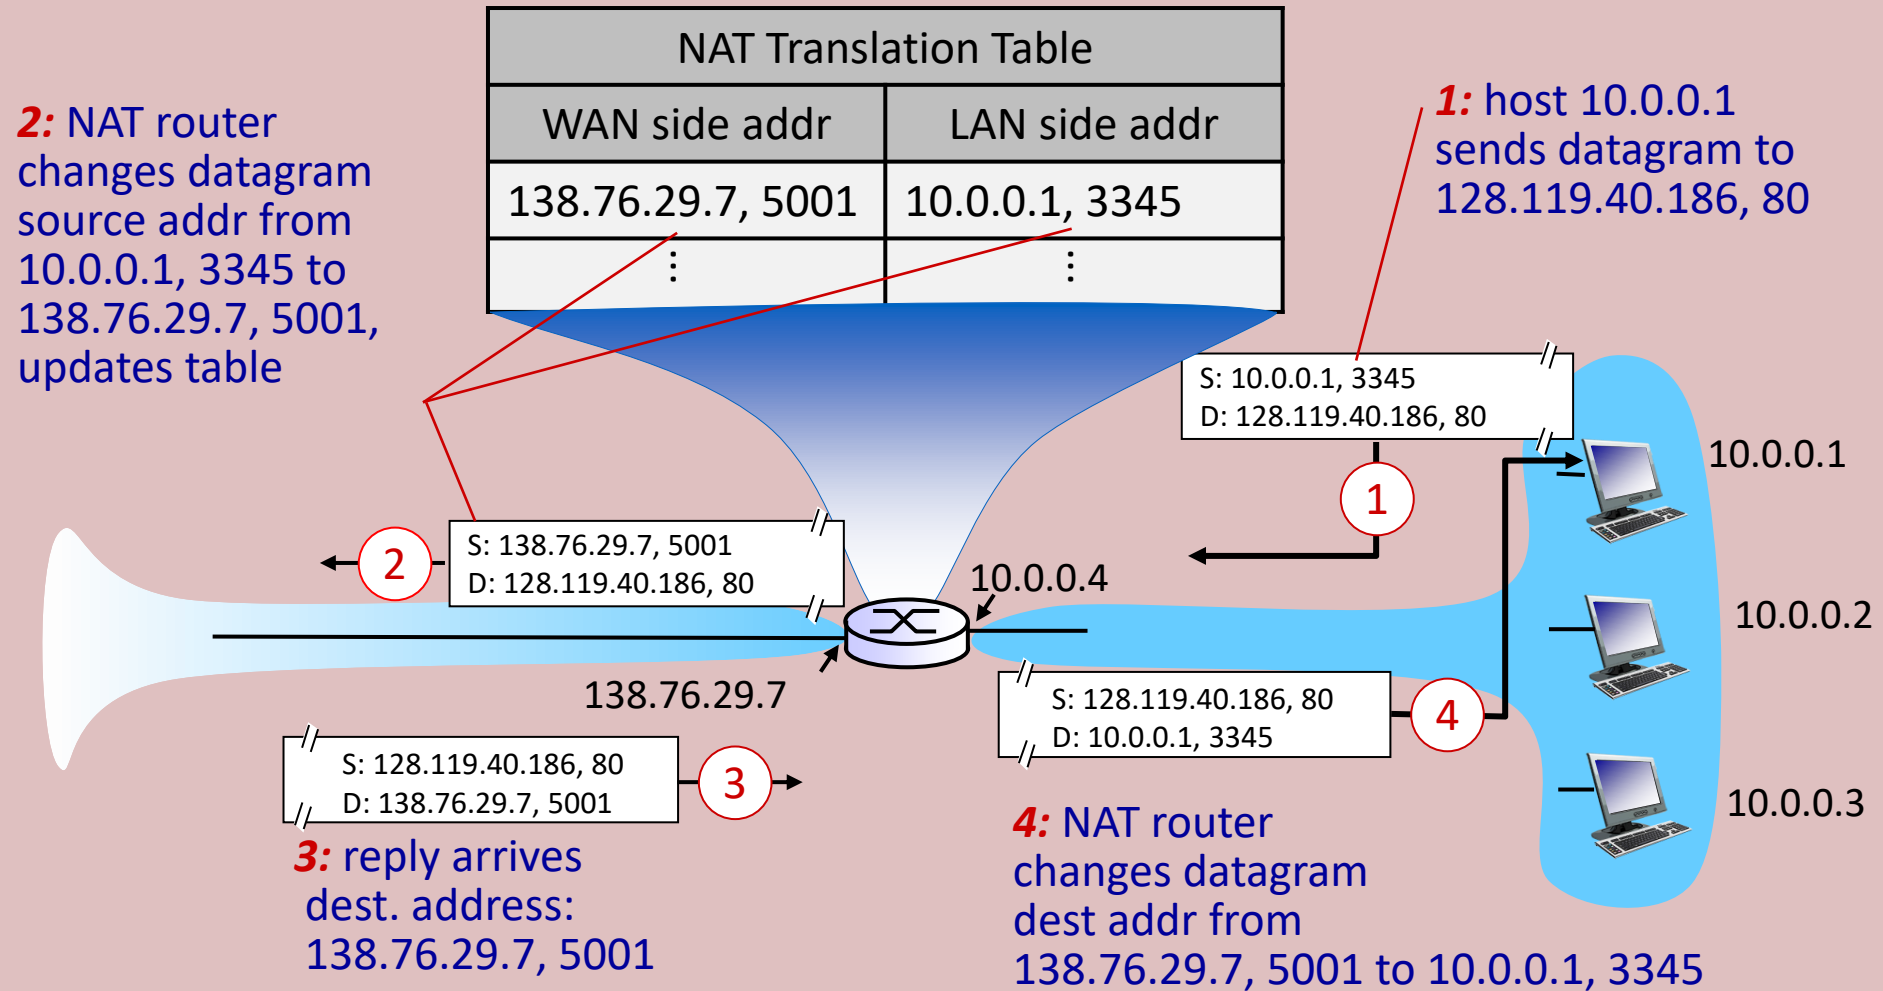
# NAT: Motivation

- Only need a single IP address from ISP
- Can change IP addresses of local network devices without notifying outside world
- Can change ISPs without changing IP addresses of local network devices
- Security bonus: local network devices not explicitly addressable by outside world

# NAT: Implementation

- Translation table: map public IP address/port to private address/port
- Outgoing packets: translate private address/port to public address/port
- Incoming packets: translate public address/port to private address/port

# NAT





# NAT

- 16-bit port number field
  - How many devices can be on a private network?
- Used to be somewhat controversial
  - Routers should strictly be layer 3
  - IPv6 should solve address shortage
  - Violates end-to-end argument
    - App designers must account for NAT
    - E.g., P2P applications like Skype
- NAT traversal: clients connecting to servers behind NATs

# IPv6

- Replaces 32-bit IP addresses with 128-bit IP addresses (why?)
- New header format to speed up packet processing/forwarding
- Header changes to facilitate QoS
- Fixed-length 40-byte header
- No fragmentation

# IPv6 Datagram Header

version (4 bits)	diffserv (4 bits)	ECN (2)	flow label (20 bits)		
payload length (16 bits)			next header (8 bits)	hop limit (8 bits)	
source IP address (128 bits)					
destination IP address (128 bits)					

# IPv6 Packet Header Fields

- Flow label: identify packets in same flow
- Next header: identify upper layer protocol for payload data
- Hop limit: similar to TTL


# Other changes from IPv4

- Checksum: removed entirely to reduce processing time at each hop
- Options: allowed, but outside of header, indicated by “Next Header” field
- ICMPv6: new version with multicast group management

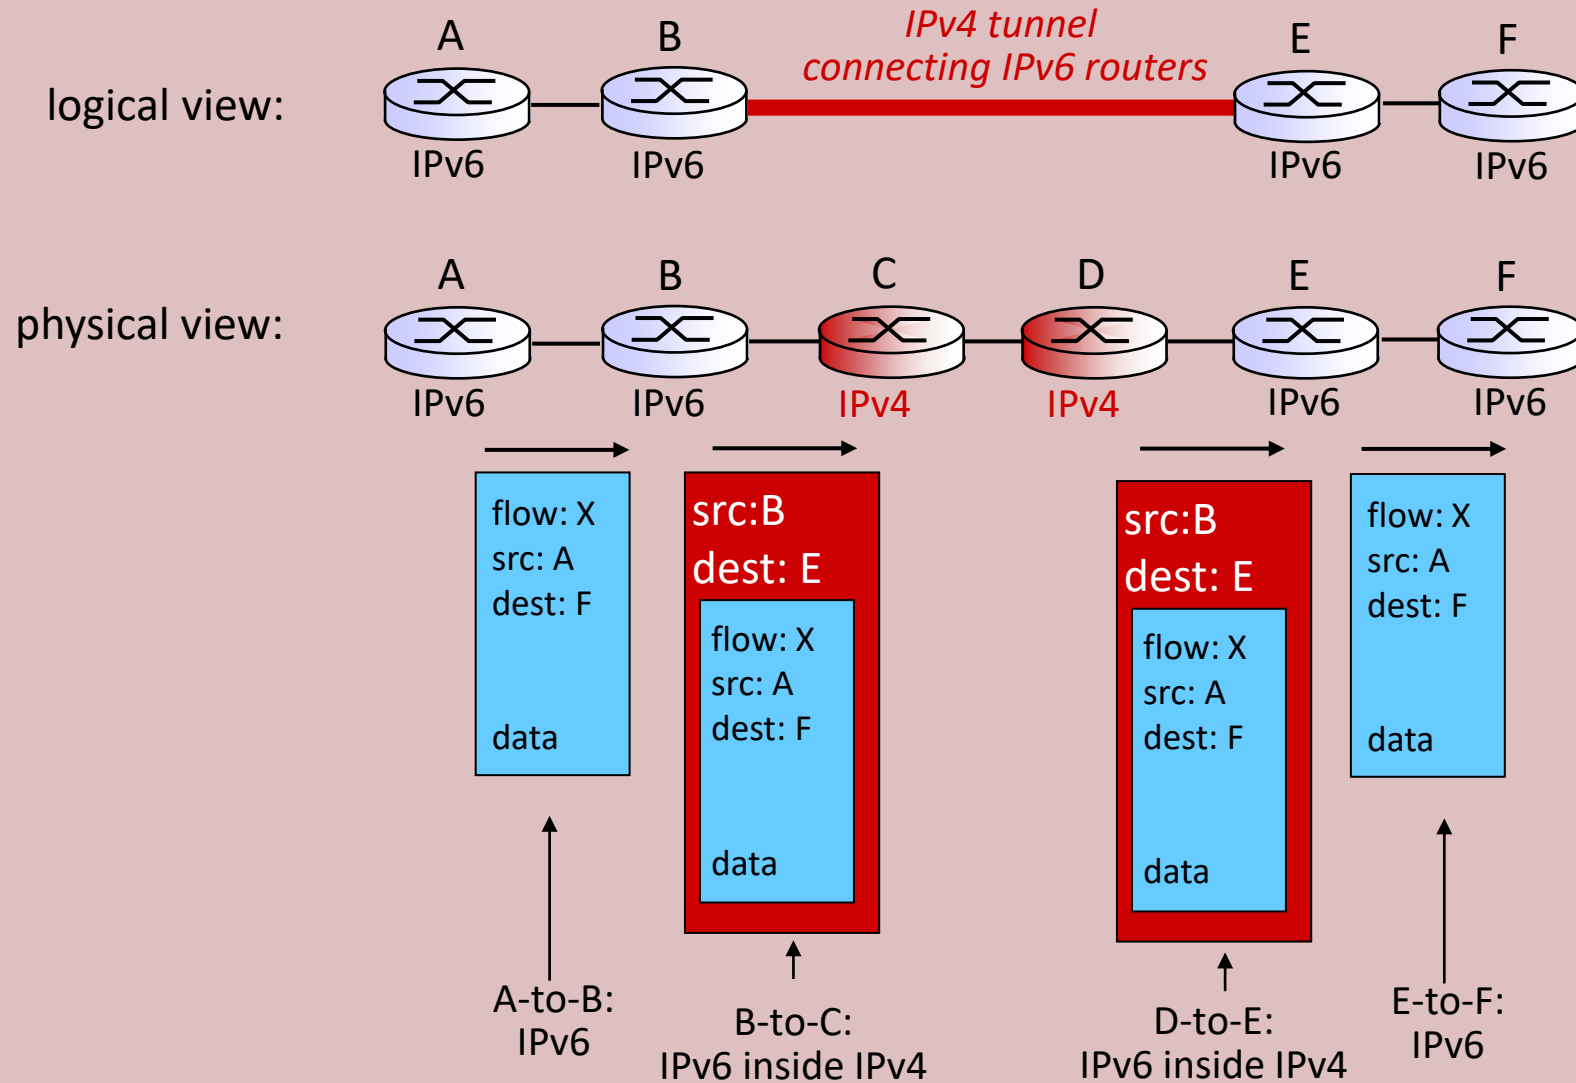
# IP Tunneling

- Routers cannot all be upgraded simultaneously (i.e., no “flag days”)
- Network must operate with both IPv4 and IPv6 routers (dual stack)
- Tunneling: carry IPv6 packet as payload in IPv4 among IPv4 routers

# IP Tunneling



# IP Tunneling





# IPv6 Adoption

- Google: 12.9% of clients access services via IPv6
- NIST: One-third of all US government domains are IPv6 capable
- Deployment is 20 years and counting (why?)

# Thank You!

# Networks

Generalized Forwarding and SDN

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Generalized Forwarding and SDN

## Section 4.4

# Generalized Forwarding

- Thus far, routing tables have only considered destination-based forwarding
- Flow tables consider more than just destination IP addresses
- Logically centralized controller computes and distributes flow tables (software defined networking)

# Generalized Forwarding

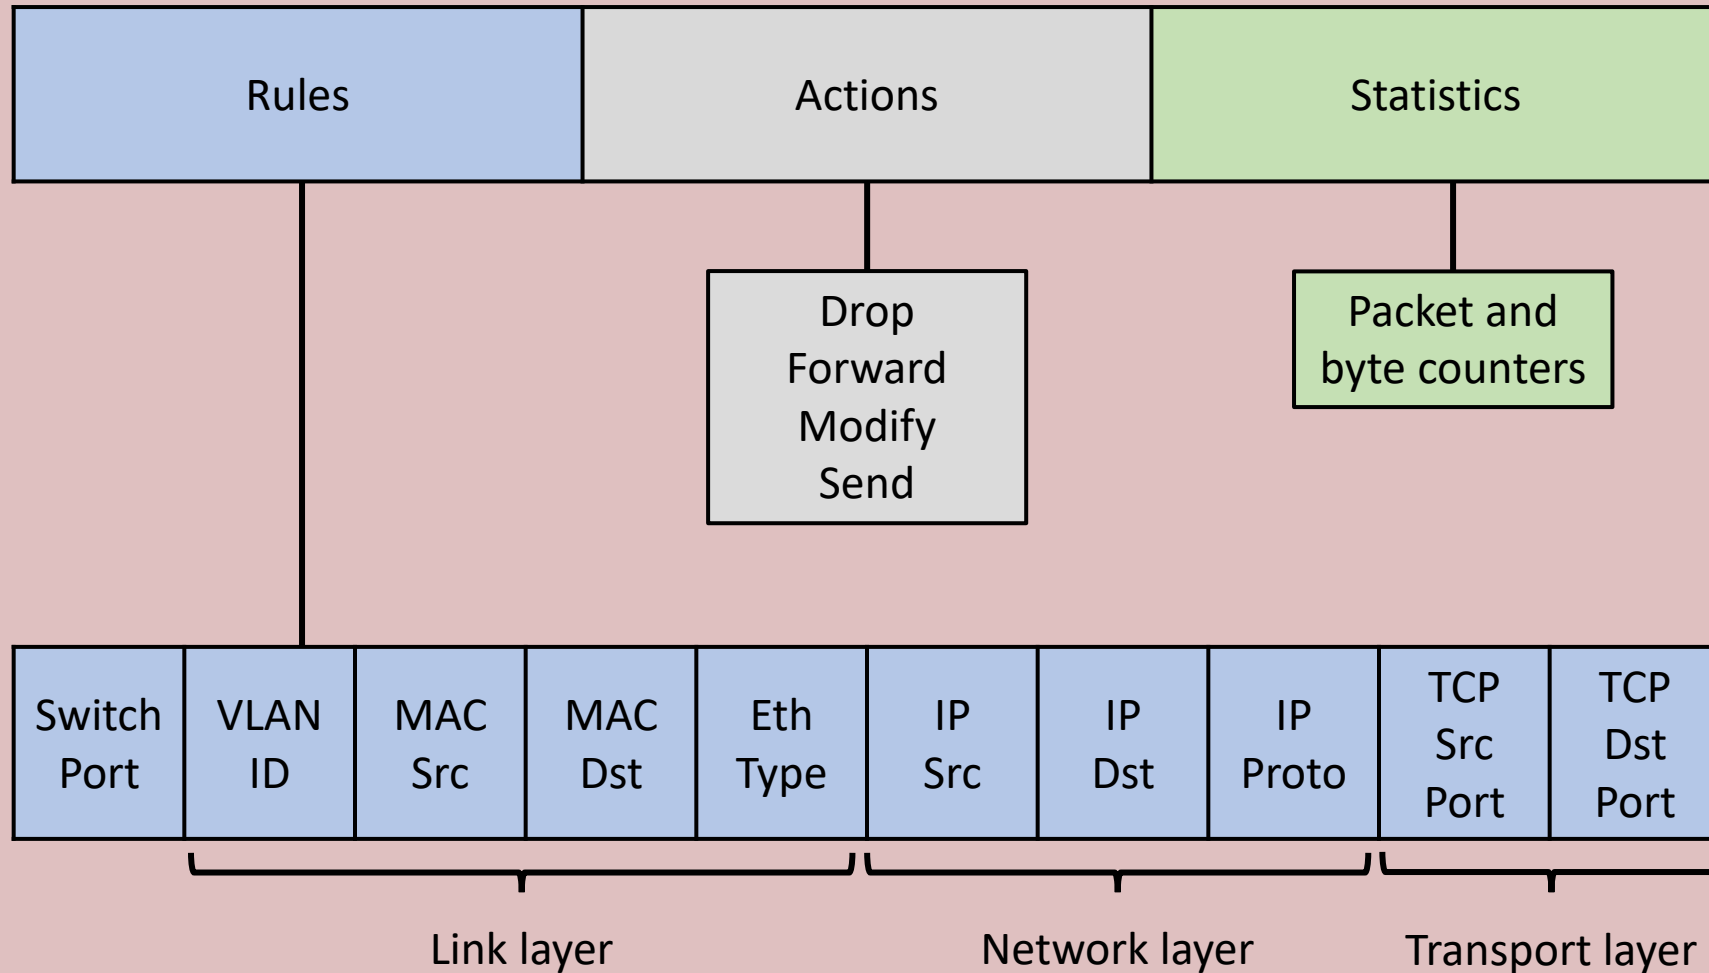
- More than just forwarding a packet to a destination (router)
- The following actions are also supported
  - Drop packet (firewall)
  - Modify packet (NAT)
  - Send packet to SDN controller

# OpenFlow Data Plane Abstraction

- Flow: defined by header fields
- Flow table: defines match + action rules for flows
- Generalized forwarding: flow-based packet-handling rules
  - Match flow patterns
  - Perform action for matched packets



# OpenFlow Table Format



# OpenFlow Table Examples

Destination-Based Forwarding: datagrams headed towards 128.10.11.12 should be forwarded on output port 6

IP Src	IP Dst	Action
*	128.10.11.12	forward(6)

# OpenFlow Table Examples

Firewall: block all datagrams headed towards TCP port 22 (SSH traffic)

IP Src	IP Dst	TCP Src Port	TCP Dst Port	Action
*	*	*	22	drop

# OpenFlow Table Examples

Firewall: block all datagrams sent by 128.10.11.12

IP Src	IP Dst	Action
128.10.11.12	*	drop

# OpenFlow Table Examples

Layer-2 Switch: forward all frames from AA:BB:CC:DD:EE:FF on output port 3

MAC Src	MAC Dst	Action
AA:BB:CC:DD:EE:FF	*	forward(3)

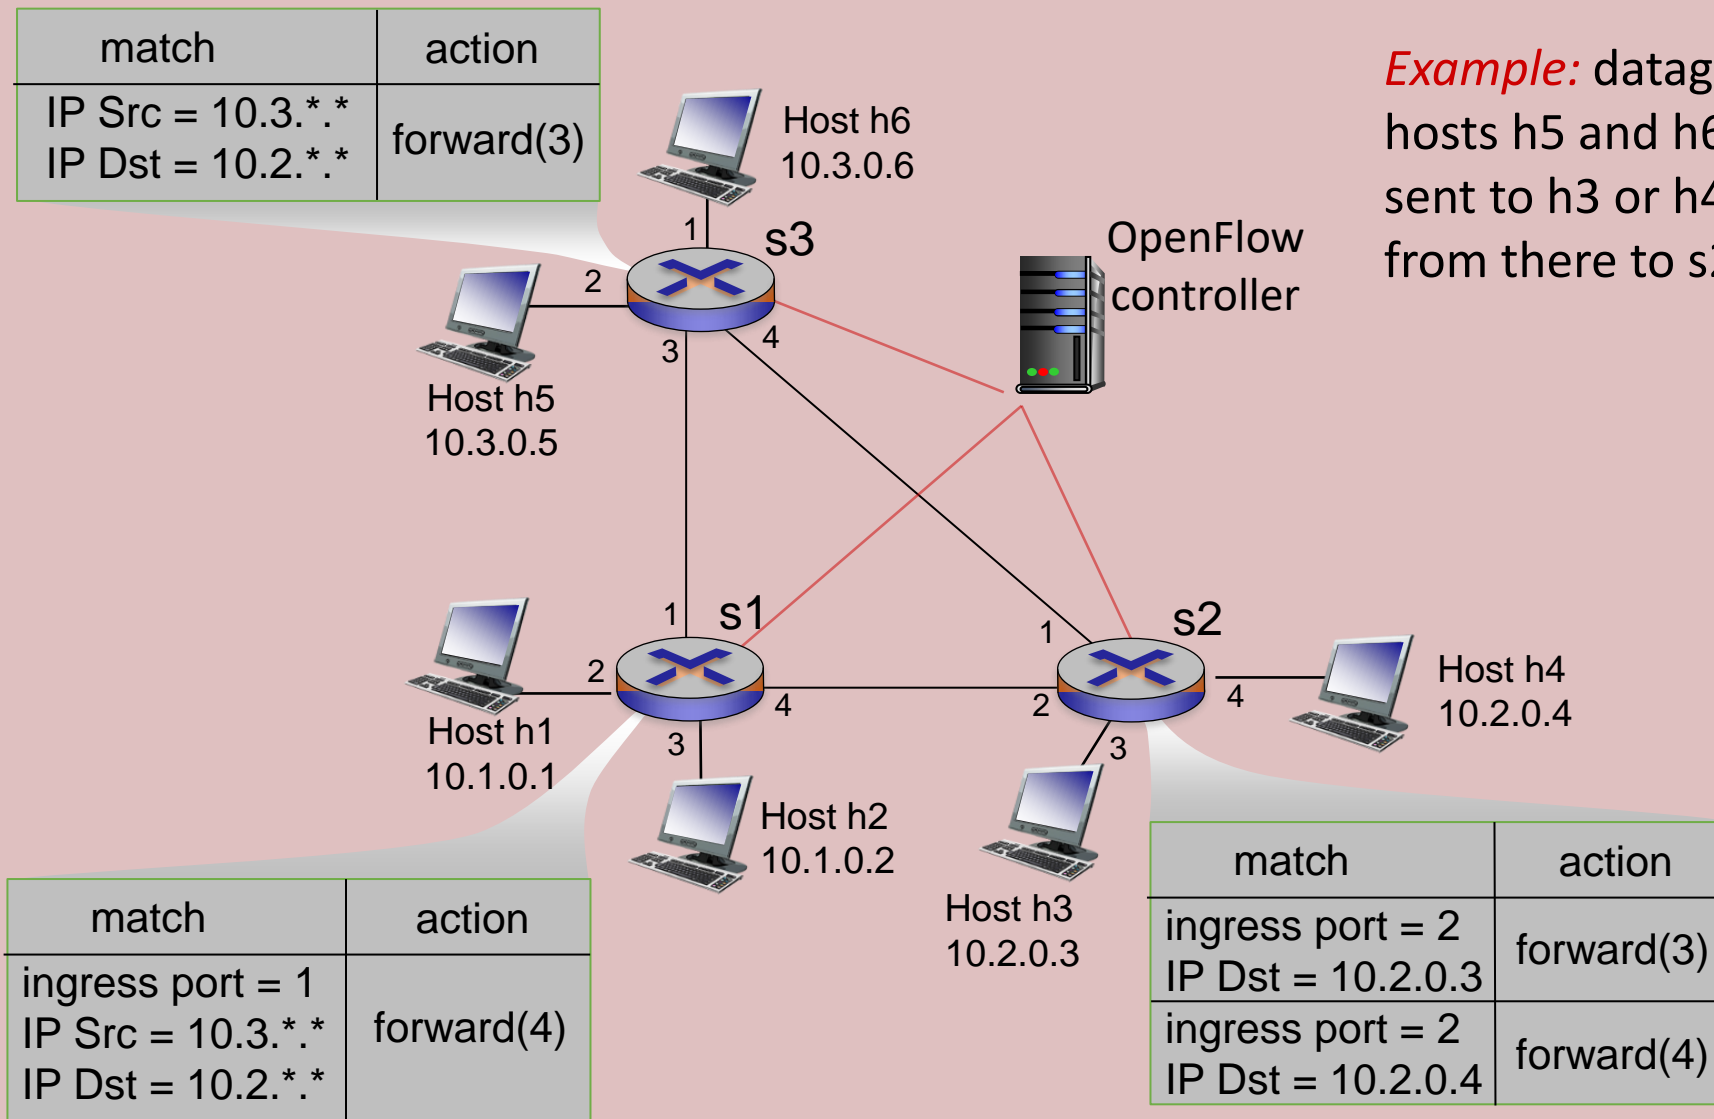
# OpenFlow Abstraction

- match+action: paradigm unifies different kinds of devices into one rule set
- Router
  - match: longest destination IP prefix
  - action: forward out on link
- Switch
  - match: destination MAC address
  - action: forward or flood

# OpenFlow Abstraction

- match+action: paradigm unifies different kinds of devices into one rule set
- Firewall
  - match: IP address and port
  - action: permit or deny
- NAT
  - match: IP address and port
  - action: rewrite address and port

# OpenFlow Example





# Thank You!