

# Networks

## Control Plane - Introduction

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Control Plane - Introduction

## Section 5.1


# Review: Network Layer Functions

- Forwarding (data plane): move packets from router's input ports to appropriate output ports
- Routing (control plane): determine route taken by packets from source to destination

# Control Plane Approaches


- Distributed per-router control (traditional routing)
  - Link state (global view)
  - Distance vector (decentralized view)
- Logically centralized control (software defined networking)

# Per-Router Control Plane



Each router participates in traditional *routing algorithm* to compute its forwarding table

# Logically Centralized Control Plane



Remote SDN controller computes and distributes forwarding tables.  
Routers communicate to controller via controller agents.

# Vacation Analogy

- Group of friends collectively figure out daily sightseeing schedule during breakfast (traditional routing algorithm)
- Travel agent emails daily sightseeing schedule to group each morning (software-defined networking)



# Thank You!

# Networks

## Routing Algorithms

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Routing Algorithms

## Section 5.2

# Routing Protocols


- Objective: determine “good” paths (i.e., routes) from source to destination
  - Path: sequence of routers that packets will traverse from source host to destination host
  - “Good”: determined by some metric (e.g., least cost, fastest, least congested)

# Network Graph Abstraction

Graph  $G = (N, E)$

$N$  = set of routers =  $\{u, v, w, x, y, z\}$


$E$  = set of links =  $\{(u,v), (u,w), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z)\}$



# Network Graph Abstraction

$c(x, x') = \text{cost of link } (x, x')$

$c(w, z) = 5$



Cost could always be 1, or inversely related to bandwidth, or related to congestion

Cost of path  $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

**Key question:** what is the least cost path between u and z?

**Routing algorithm:** algorithm that finds the least cost path

# Routing Algorithm Classification

Global View	<ul style="list-style-type: none"><li>• All routers have complete topology and link cost information</li><li>• Exchange information via broadcast and/or flooding</li><li>• <i>Link state algorithms</i></li></ul>
Decentralized View	<ul style="list-style-type: none"><li>• Routers only know immediate neighbors and link cost to neighbors</li><li>• Exchange information with neighbors</li><li>• Iterative computation</li><li>• <i>Distance vector algorithms</i></li></ul>

Static	<ul style="list-style-type: none"><li>• Routes change slowly over time</li><li>• Might involve manual configuration</li></ul>
Dynamic	<ul style="list-style-type: none"><li>• Routes change relatively fast</li><li>• Triggered by periodic updates</li><li>• Triggered by link cost changes</li></ul>



# Dijkstra's Algorithm

- Link state routing algorithm
- Network topology and link costs known by all nodes via **link state broadcasts**
- Compute least cost paths from single source node to all other nodes
- After  $k$  iterations, know least cost path to  $k$  destinations

# Dijkstra's Algorithm

$c(x,y)$	Link cost from node $x$ to $y$ ; set to $\infty$ if not direct neighbors
$D(v)$	current value of cost of path from single source to destination $v$
$p(v)$	predecessor node along path from source to $v$
$N'$	set of nodes whose least cost path is definitively known

# Dijkstra's Algorithm - Initialization

$N' = \{u\}$

for all nodes  $v \in N$

if  $v$  adjacent to  $u$

$D(v) = c(u, v)$  // Link cost between  $u$  and  $v$

else

$D(v) = \infty$

# Dijkstra's Algorithm – Main Loop

While some nodes are not in  $N'$ :

Find  $w$  not in  $N'$  such that  $D(w)$  is a minimum

Add  $w$  to  $N'$

For all  $v$  adjacent to  $w$  and not in  $N'$

// Update cost if known shortest path to  $w$  plus cost from  $w$  to  $v$  is shorter than current cost


$D(v) = \min(D(v), D(w) + c(w, v))$

Also, if new path,  $p(v) = w$

# Example: Dijkstra's Algorithm

Step	N'	D(v), p(v)	D(w), p(w)	D(x), p(x)	D(y), p(y)	D(z), p(z)
0	u	7, u	3, u	5, u	$\infty$	$\infty$
1	uw	6, w		5, u	11, w	$\infty$
2	uwx	6, w			11, w	14, x
3	uwxv				10, v	14, x
4	uwxvy					12, y
5	uwxvyz					

Shortest path tree from source to destinations can be constructed by tracing predecessor nodes



# Dijkstra's Algorithm Complexity

$N$  = set of nodes

$n = |N|$ , or the number of nodes in  $N$

Checks all nodes,  $w$ , not in  $N'$  on each iteration

$n(n+1)/2$  comparisons:  $O(n^2)$

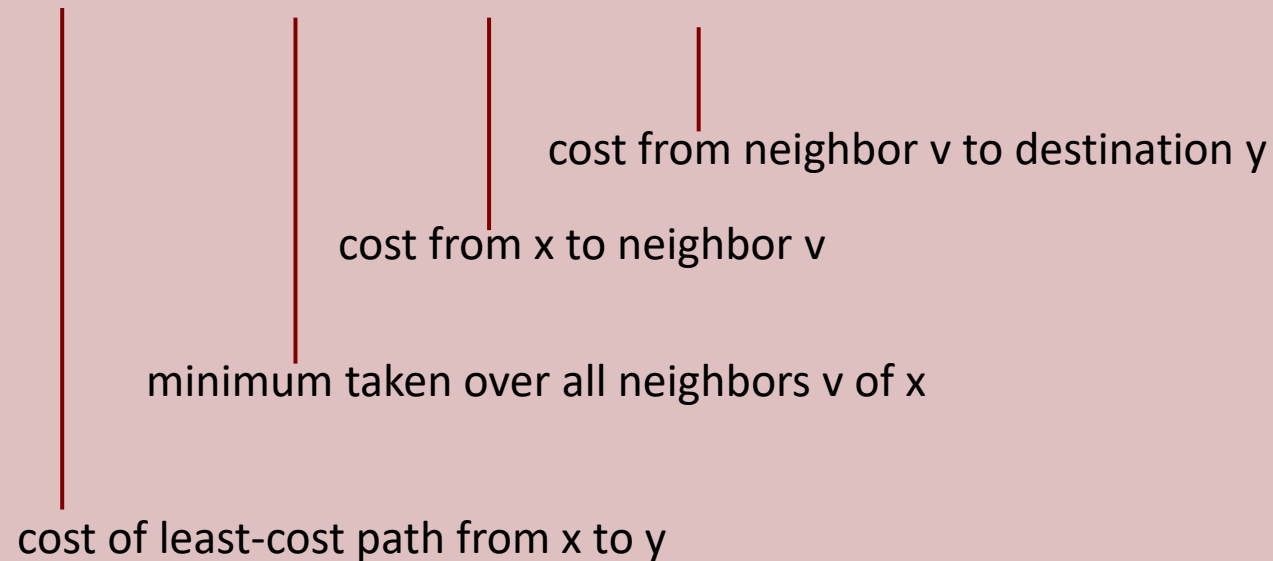
More efficient implementations are possible:  $O(n \log n)$

# Bellman-Ford Algorithm

- Distance vector routing algorithm
- Each router knows distance to neighbors, but *has no global view of network topology*
- Neighbors exchange **distance vectors** containing their best known distances to each destination

# Bellman-Ford Equation

$$d_x(y) = \min_v \{ c(x, v) + d_v(y) \}$$






# Example: Bellman-Ford

$$d_v(z)=5, d_x(z)=3, d_w(z)=3$$

Bellman-Ford equation says:

$$\begin{aligned} d_u(z) &= \min\{c(u,v)+d_v(z), c(u,x)+d_x(z), c(u,w)+d_w(z)\} \\ &= \min\{2 + 5, 1 + 3, 5 + 3\} \\ &= 4 \end{aligned}$$



What's is u's next hop to z?

The node achieving the minimum is the next hop in the shortest path, used in the forwarding table

# Distance Vector Algorithm

Wait for local link cost change or distance vector update from neighbor

Recompute local distance vector using Bellman-Ford equation

Notify neighbors of updated distance vector if any entries change

Repeat


# Example: Distance Vector

t=0 (x only knows local links)

$DV_x = (0, 2, 7)$

$DV_y = (\infty, \infty, \infty)$

$DV_z = (\infty, \infty, \infty)$



t=1 (x exchanges with y)

$DV_x = (0, 2, 3)$

$DV_y = (2, 0, 1)$

$DV_z = (\infty, \infty, \infty)$

Better distance found to  
z, so notify neighbors of  
latest  $DV_x$

t=2 (x exchanges with z)

$DV_x = (0, 2, 3)$

$DV_y = (2, 0, 1)$

$DV_z = (7, 1, 0)$

No change, so do not  
send any updates


# Link Cost Improvements

t=0

$DV_x = (0, 4, 5)$

$DV_y = (4, 0, 1)$

$DV_z = (5, 1, 0)$



t=1 (link cost improvement)

$DV_x = (0, 1, 2)$

$DV_y = (1, 0, 1)$

$DV_z = (5, 1, 0)$

x and y update distance vectors, then notify neighbors

t=2 (exchange with z)

$DV_x = (0, 1, 2)$

$DV_y = (1, 0, 1)$

$DV_z = (2, 1, 0)$

z updates distance vector and notifies neighbors, but no further updates

# Count-to-Infinity Problem

t=0

$$d_y(x) = \min(4, 1+5) = 4$$

$$d_z(x) = \min(50, 1+4) = 5$$

t=1 (link cost increase)

$$d_y(x) = \min(60, 1+5) = 6$$

$$d_z(x) = \min(50, 1+6) = 7$$

t=2


$$d_y(x) = \min(60, 1+7) = 8$$

$$d_z(x) = \min(50, 1+8) = 9$$

t=3

$$d_y(x) = \min(60, 1+9) = 10$$

$$d_z(x) = \min(50, 1+10) = 11$$



$$d_y(x) = \min\{ c(y,x), c(y,z) + d_z(x) \}$$
$$d_z(x) = \min\{ c(z,x), c(z,y) + d_y(x) \}$$

When will the distance vectors  
finally converge to stable  
values?

# Poison Reverse

- Count-to-infinity is fundamentally caused by circular dependency between  $z$  and  $y$
- Breaking this cycle should solve the issue
- If some node  $u$  depends on node  $v$  to reach some destination  $w$ , then  $u$  advertises *infinite distance* to  $w$  to  $v$

# Poison Reverse

t=0

$$d_y(x) = \min(4, 1+\infty) = 4$$

$$d_z(x) = \min(50, 1+4) = 5$$

t=1 (link cost increase)


$$d_y(x) = \min(60, 1+\infty) = 60$$

$$d_z(x) = \min(50, 1+60) = 50$$

t=2

$$d_y(x) = \min(60, 1+50) = 51$$

$$d_z(x) = \min(50, 1+\infty) = 50$$



$$d_y(x) = \min\{ c(y,x), c(y,z) + d_z(x) \}$$
$$d_z(x) = \min\{ c(z,x), c(z,y) + d_y(x) \}$$

# Thank You!



# Networks

## Intra-AS Routing

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Intra-AS Routing

## Section 5.3

# Autonomous Systems

- Internet = network of networks
- Network admins want to control routing in their own networks
- Aggregate routers into **autonomous systems (ASes)** under single administrative domains

# Intra-AS Routing

- Routing amongst hosts and routers within the same AS
- All routers in the same AS must run the same intra-domain protocol
- Routers in different ASes can run different intra-domain protocols
- Gateway router: “edge” of its own AS with links to routers in other ASes
- Also known as interior gateway protocols (IGP)

# Intra-AS Routing

- Most common intra-AS routing protocols
  - RIP: Routing Information Protocol
  - OSPF: Open Shortest Path First
  - IGRP: Interior Gateway Routing Protocol (Cisco proprietary protocol until 2016)

# Open Shortest Path First (OSPF)

- “Open”: publicly available (RFC 2328)
- Link-state algorithm
- Messages sent directly over IP
- Routers flood link-state advertisements for each attached link to all other routers in the entire AS

# OSPF Features

- All messages are authenticated
- Multiple same-cost paths are allowed (only one path in RIP)
- Integrated multicast support by simply reusing shortest path tree
- Hierarchical OSPF in large domains




# Hierarchical OSPF

- Two-level hierarchy: **local areas** and **backbone**
- Link-state advertisements only in local area
- Each node has detailed local area topology; only know direction (shortest path) to other areas

# Hierarchical OSPF

- **Area border routers:** “summarize” distances to networks in own area, advertise to other areas
- **Backbone routers:** run OSPF routing limited to backbone
- **Boundary routers:** connect to other ASes

# Hierarchical OSPF



# Routing Information Protocol (RIP)

- One of the first intra-AS routing protocols in TCP/IP (first defined in RFC 1058)
- Distance vector algorithm
- Uses fewer resources than OSPF (no flooding)
- Three versions: RIP-1 and RIP-2 for IPv4 and RIPng for IPv6

# RIP Updates

- Routers can pull routing information (e.g., immediately after booting up)
- Routers can push routing information
- Update timer typically fires every 30 seconds to trigger push
- Routes have expiration times (updates refresh the expiration timer)

# RIP Issues

- Slow convergence since topology changes might take minutes to converge
- Routing loops can occur in rare instances
- Counting to infinity
- Newer OSPF is supposed to replace RIP

# Thank You!

# Networks

## Inter-AS Routing



Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# Inter-AS Routing

## Section 5.4

# Inter-AS Routing

- Routing among ASes
- Gateway routers must run **inter-domain protocol** in addition to intra-domain protocol
- Inter-domain protocol must be the same across different domains

# Border Gateway Protocol (BGP)

- The de-facto inter-domain routing protocol
- *Glue* that binds the various ASes into the global internet
- Allows subnet to *advertise* its existence to the rest of the Internet
- Determine “good” routes to other networks based on reachability information *and policy*

# Internal and External BGP


- **External BGP (eBGP):** allows each AS to obtain subnet reachability information from neighboring ASes
- **Internal BGP (iBGP):** allows each AS to propagate reachability information to all AS-internal routers

# BGP Router

- BGP-capable routers are called **speakers**
- BGP neighbors are called **peers**
- Peers within **same AS** speak **iBGP**
- Peers **across ASes** speak **eBGP**

# eBGP and iBGP

- Gateway routers run eBGP and iBGP
- Internal routers only run iBGP



# BGP Session

- Peers establish **session** over semi-permanent TCP connections
- Peers exchange **path** information to network prefixes via **advertisements**
- ASes indicate willingness to forward packets along advertised paths

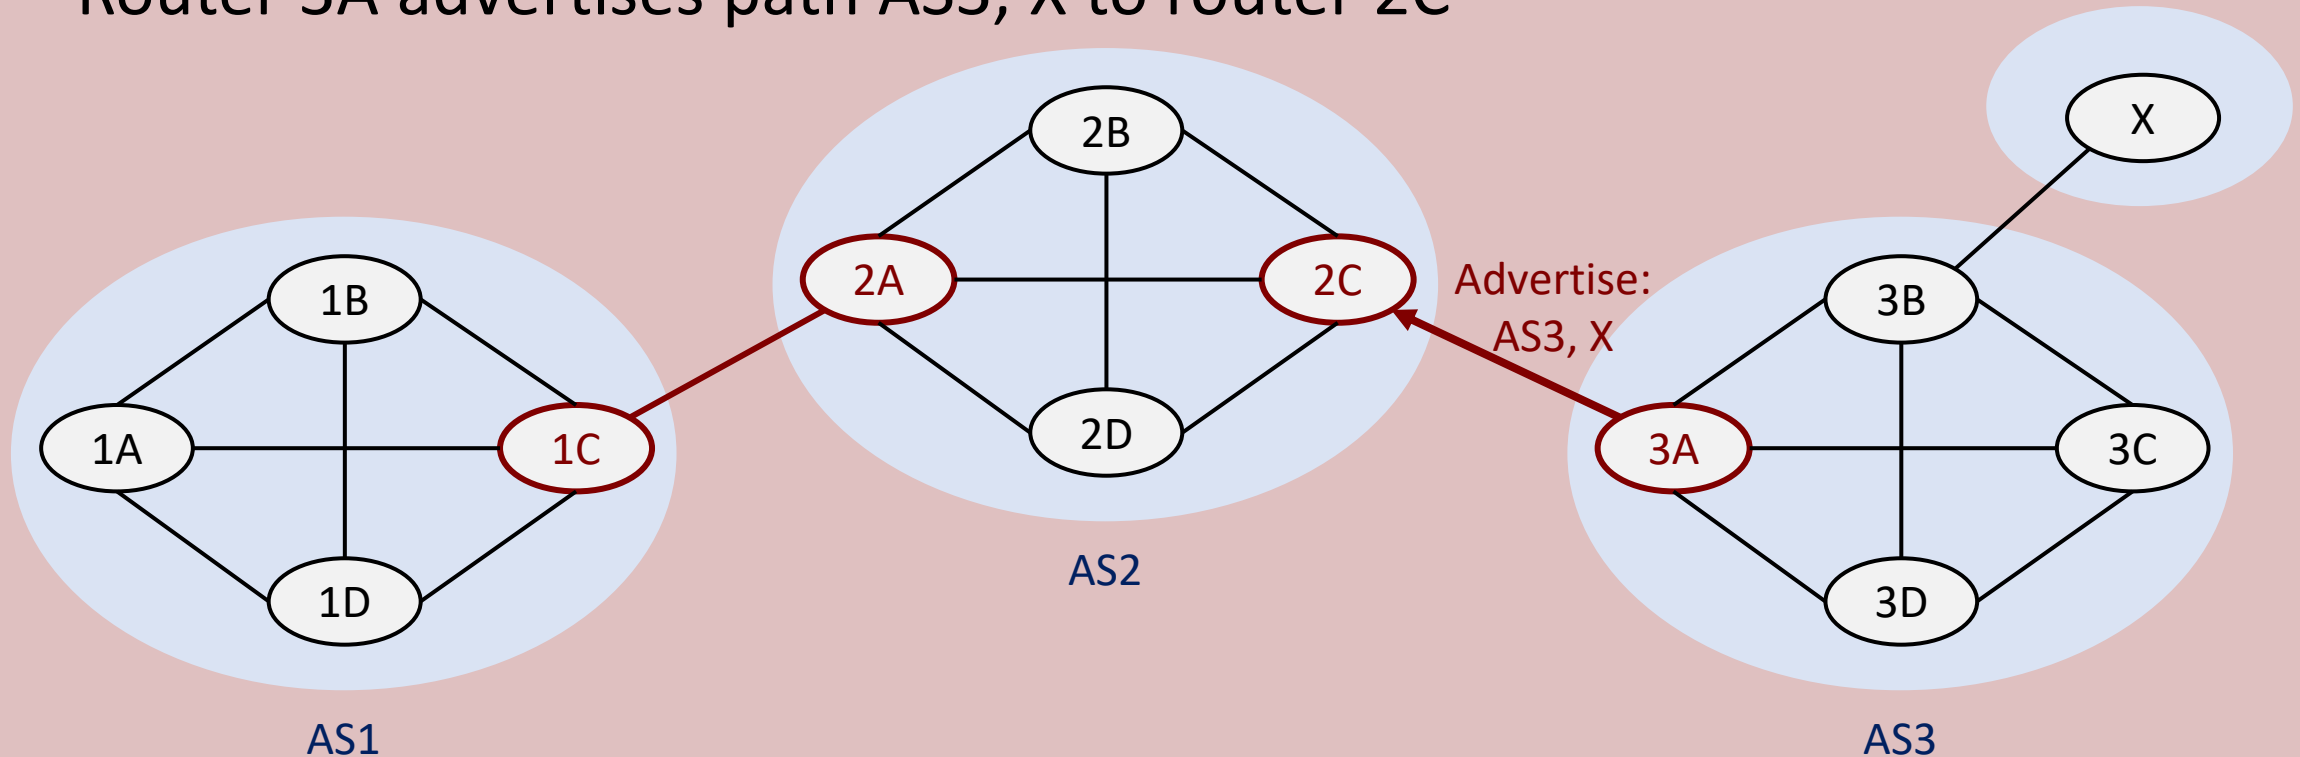


# BGP Message Types

- **OPEN** : opens TCP connection to remote BGP peer and authenticates
- **UPDATE** : advertises new path or withdraws old path
- **KEEPALIVE** : keeps connection alive in absence of UPDATES; also ACKs OPEN request
- **NOTIFICATION** : reports errors; also used to close connection

# Example: BGP Session

- AS3 is willing to route traffic from AS2 to X
- Router 3A advertises path AS3, X to router 2C



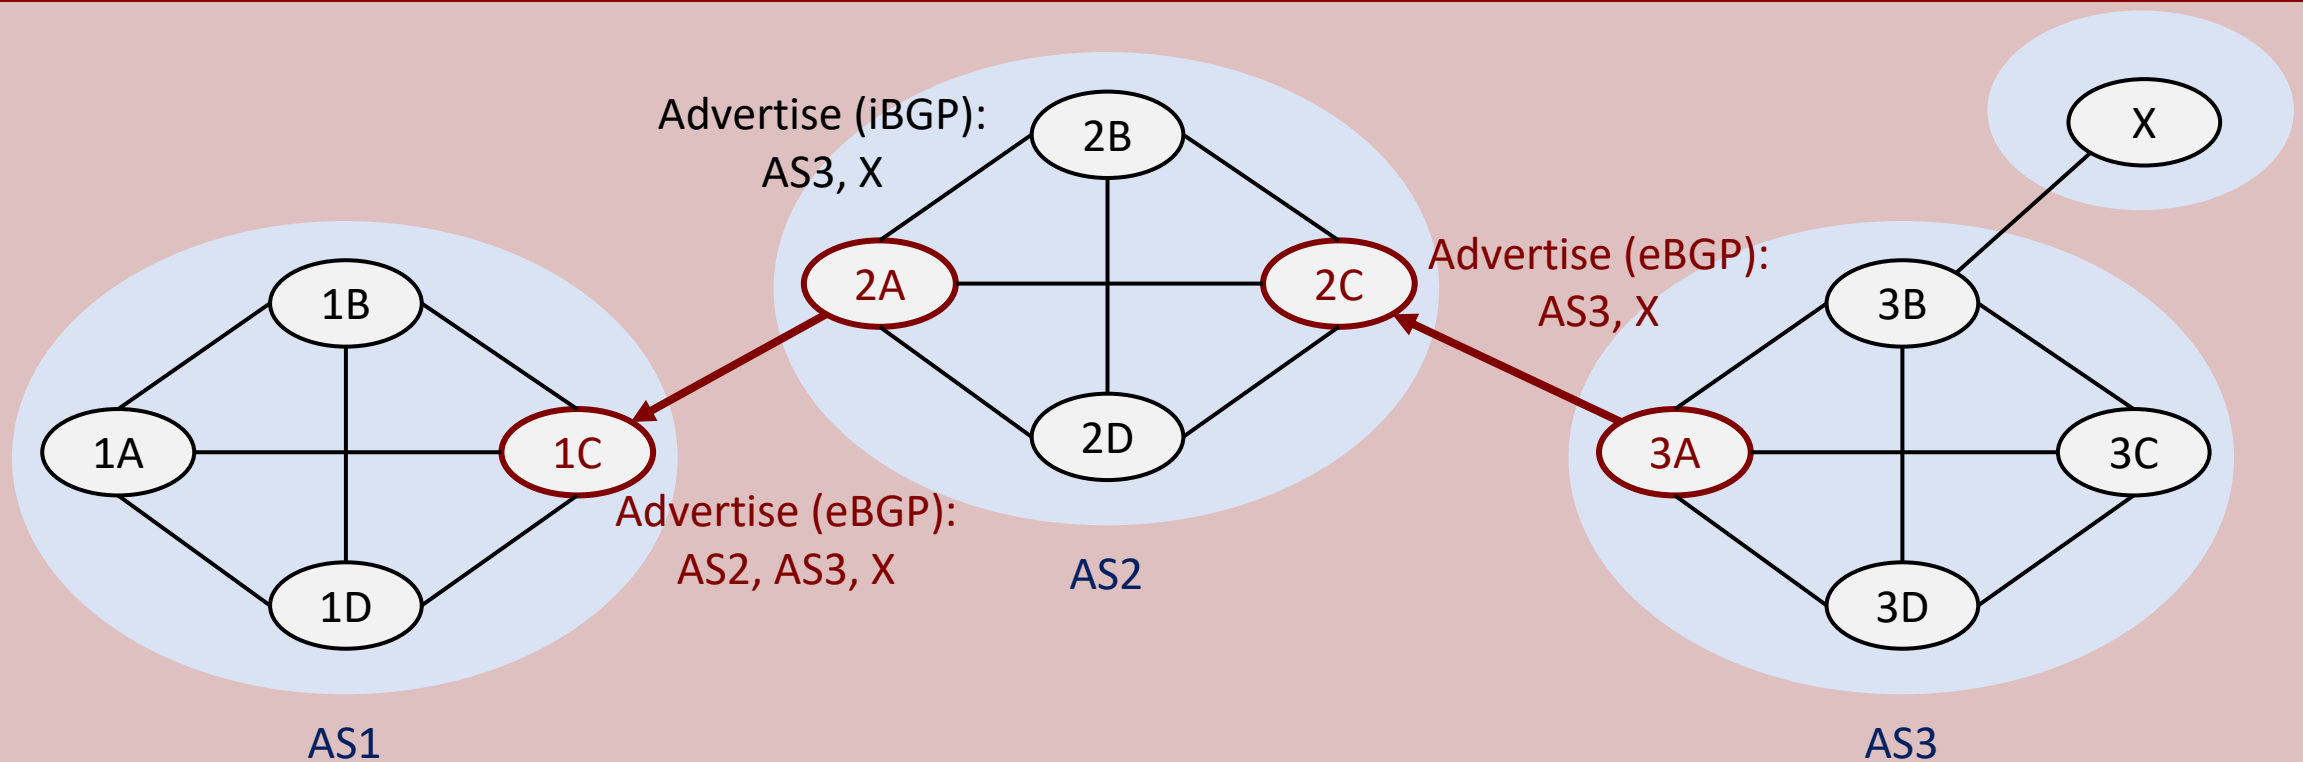
# BGP Path Attributes

- BGP considers path attributes, not just distance (unlike OSPF and RIP)
- **AS-PATH**: list of ASes through which prefix advertisement has passed
- **NEXT-HOP**: IP address router interface that begins AS-PATH (**why?**)

# BGP Policy-Based Routing

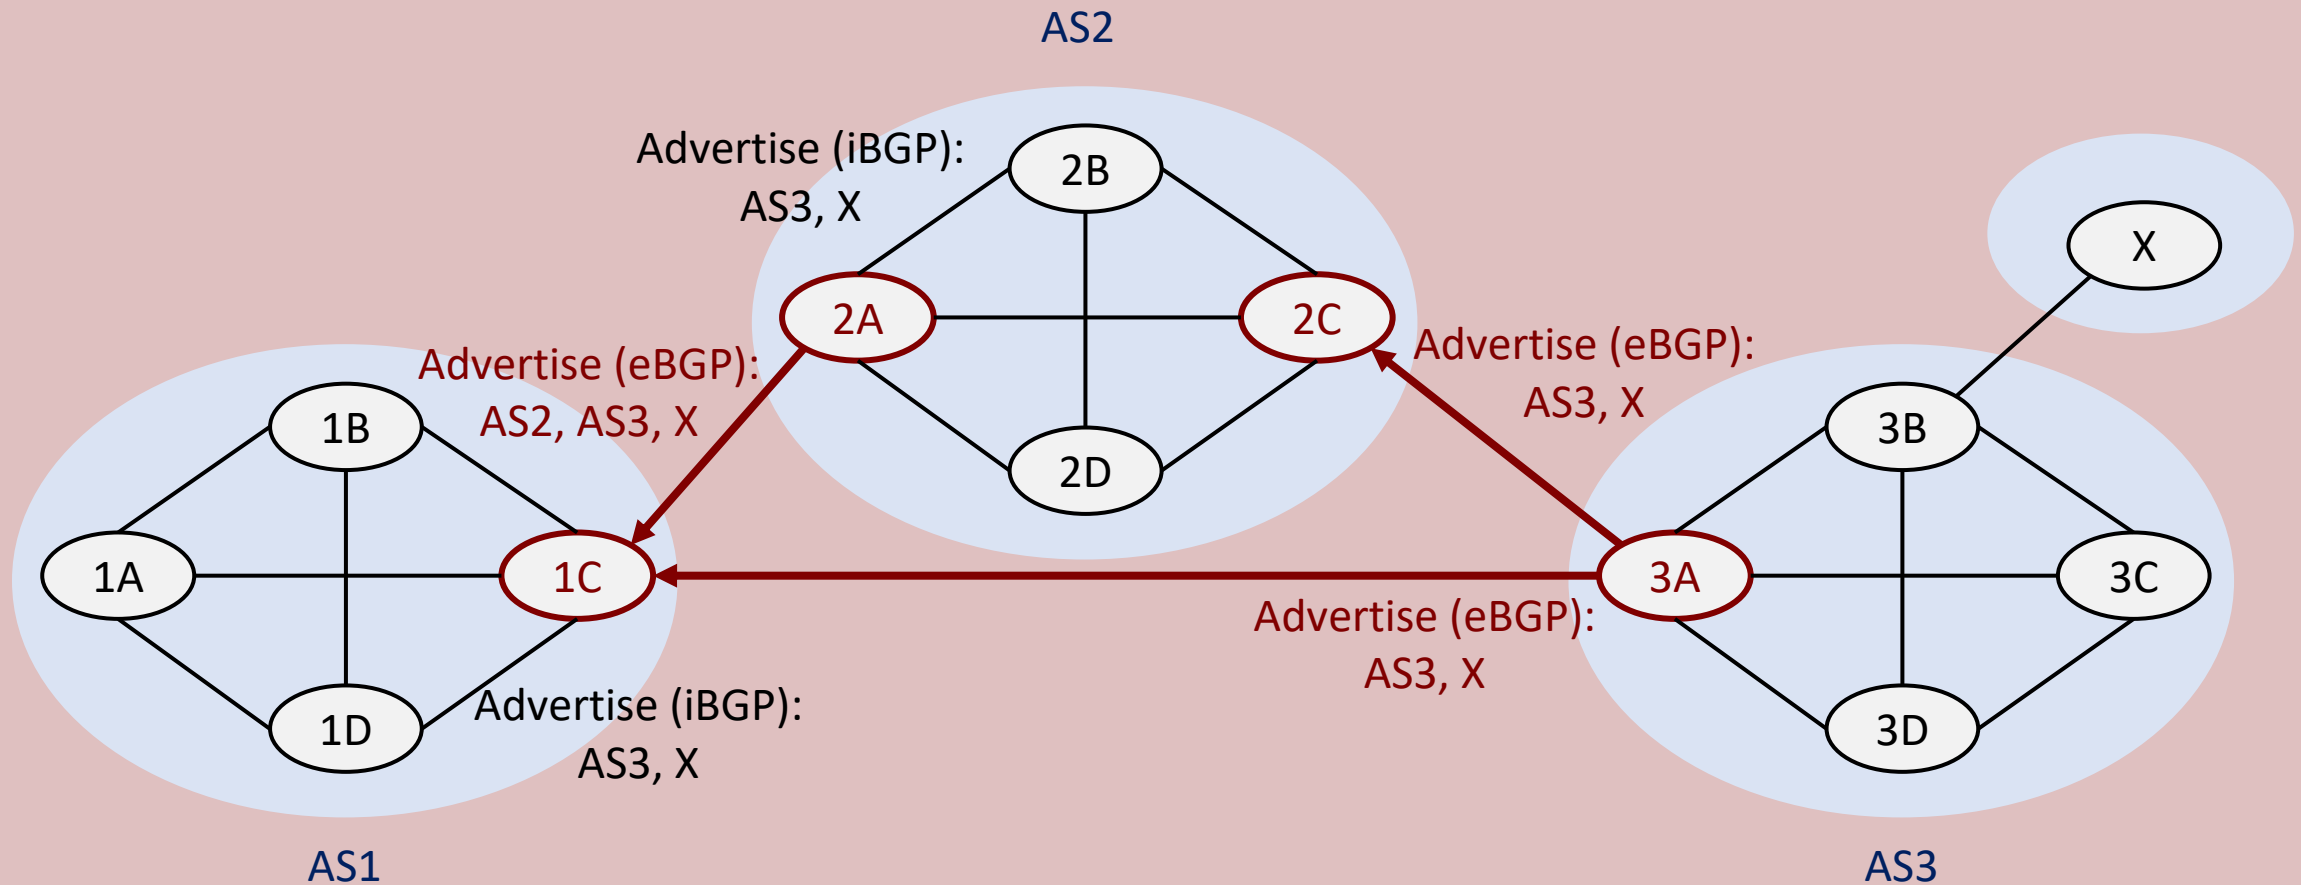
- Import policy: determines whether or not to accept path in advertisements received via gateway router (e.g., never route through AS4)
- Export policy: determines whether or not to advertise paths to other neighboring ASes (e.g., pretend that AS5 is not reachable even though it is)

# BGP Path Advertising



AS2 policy accepts path AS3, X  
AS2 policy advertises path AS2, AS3, X

# Multiple BGP Path Advertisements



AS1 policy prefers path AS3, X over AS2, AS3, X

# BGP Route Selection Criteria

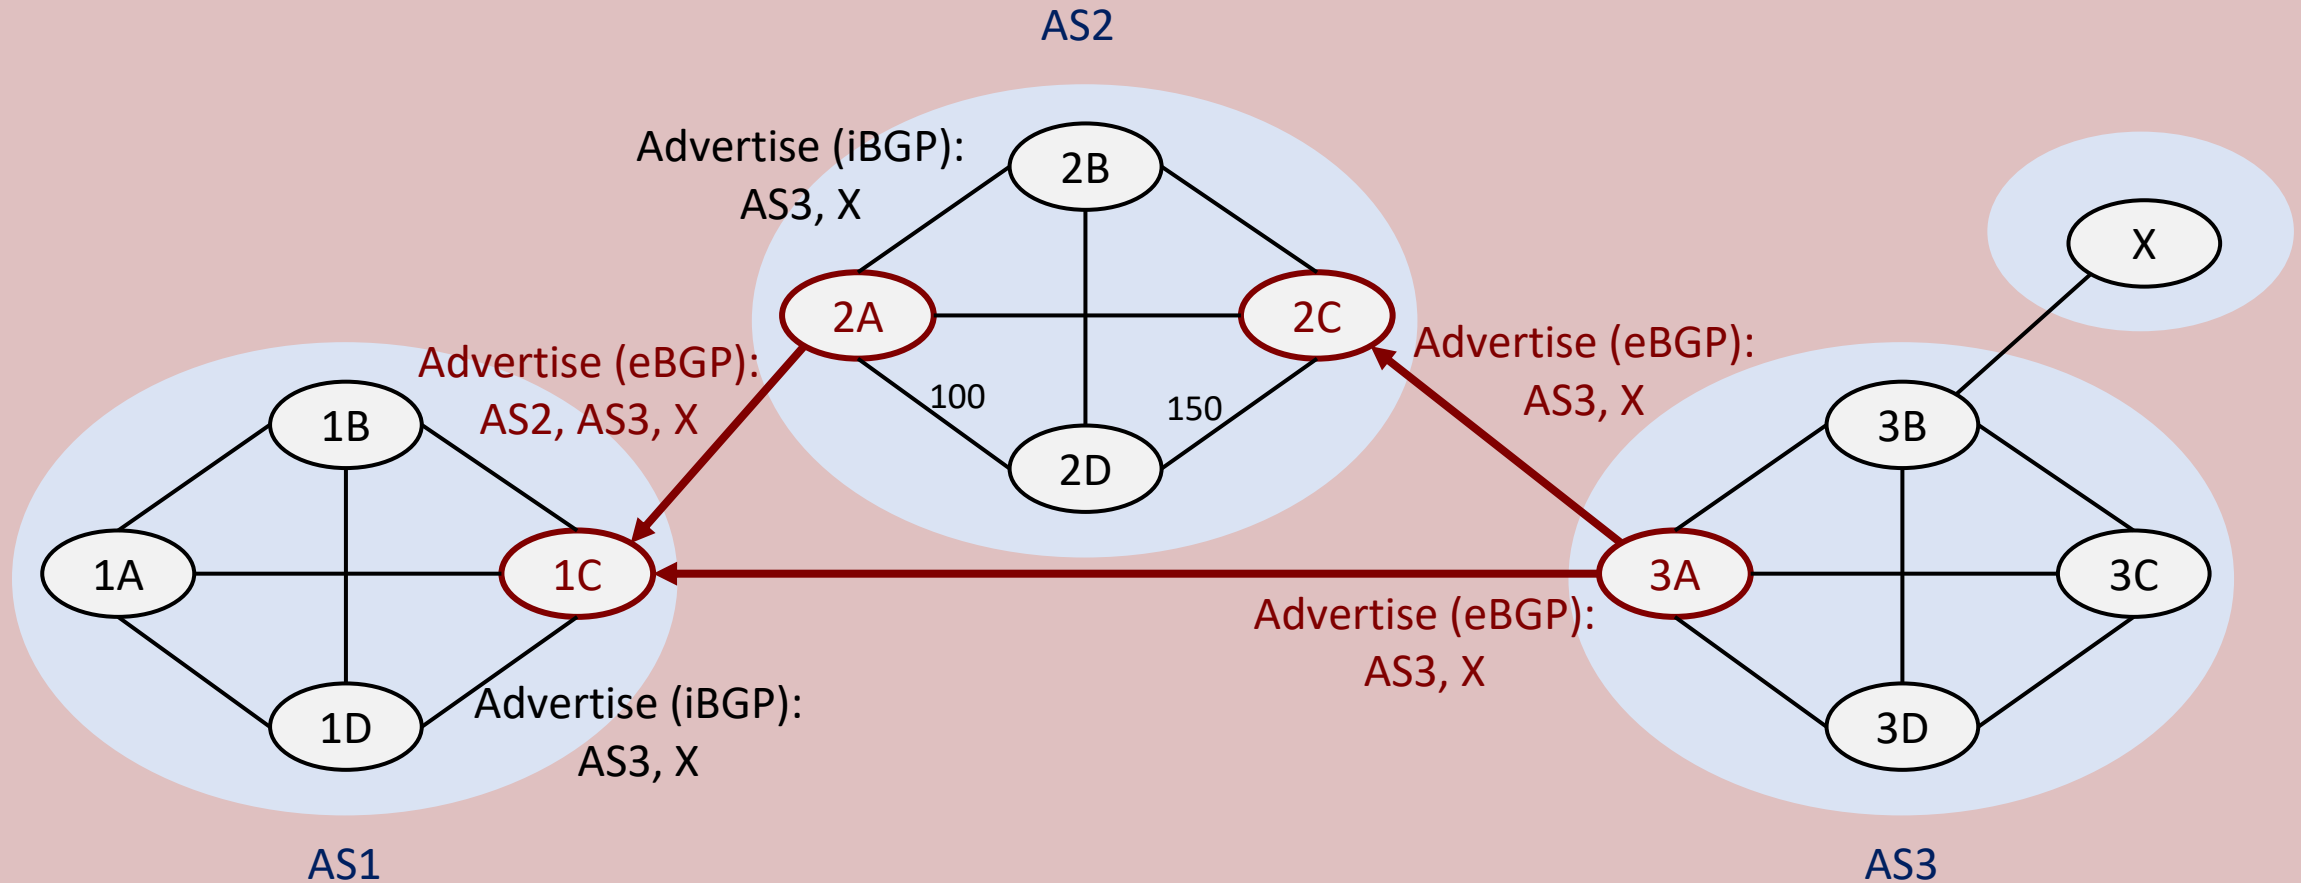
1. Local preference value attribute (**policy decision**)
2. Shortest AS-PATH
3. Closest NEXT-HOP router (**hot potato routing**)
4. Additional criteria

# Hot Potato Routing

- Optimize to achieve lowest intra-domain routing cost on your AS
- Do not worry about inter-domain routing cost for other ASes
- Can take precedence over other criteria (e.g., shortest AS-PATH)



# Hot Potato Routing



Router 2D chooses route to 2A over 2C to reach X (i.e., longer AS-PATH for lower intra-domain cost)

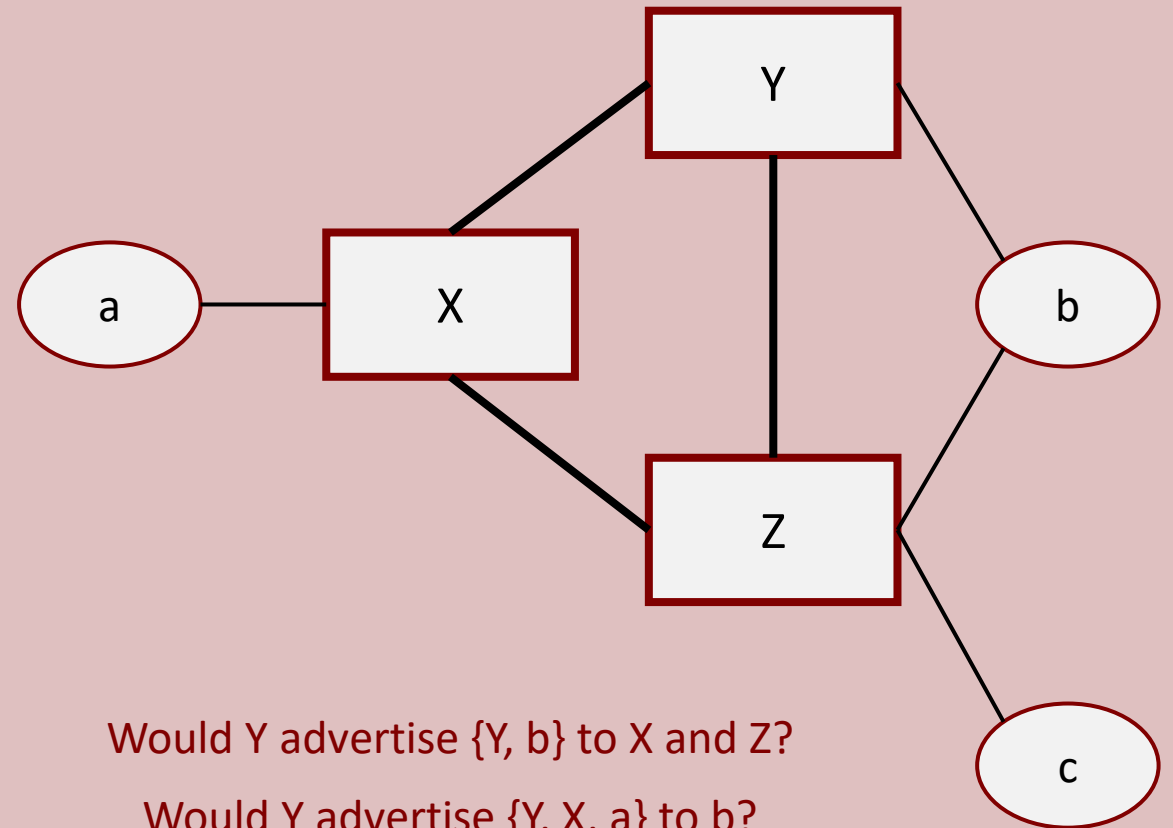
# BGP Policy-Based Advertisements

**Goal:** ISP Y only routes traffic to/from its customers

X advertises {X, a} to Y and Z

Y does not advertise {Y, X, a} to Z because it does not earn any revenue from X, Z, or a

Z routes traffic via {Z, X, a}



Would Y advertise {Y, b} to X and Z?


Would Y advertise {Y, X, a} to b?

# BGP Policy-Based Advertisements

**Goal:** Multihomed customer b does not want to transit traffic between ISPs

b suppresses its reachability to Y from Z

b suppresses its reachability to Z from Y



# Thank You!

# Networks


## SDN Control Plane

Adopted from material in “Computer Networking: A Top Down Approach” by Kurose and Ross and slides developed by William Conner

# SDN Control Plane

Section 5.5


# Review: Per-Router Control Plane



Each router participates in traditional *routing algorithm* to compute its forwarding table



# Review: Software-Defined Networking



Remote SDN controller computes and distributes forwarding tables.  
Routers communicate to controller via controller agents.

# Per-Router Control Plane

- Traditional Routing
- **Monolithic** routers running proprietary protocols on proprietary OSES (e.g., Cisco IOS)
- Separate hardware (**middleboxes**) for network functions: firewalls, load balancers, and NATs



# SDN Motivation

- Easier network management
  - Avoid router misconfigurations
  - Greater flexibility of traffic flows
- Open (non-proprietary) implementation of control plane
  - Mix-and-match generalized forwarding hardware from different vendors
- Table-based forwarding allows “programming” routers
  - Centralized “programming” is easier: compute tables centrally and distribute
  - Distributed “programming” is more difficult: compute tables as result of distributed algorithm (protocol) implemented in each and every router


# Analogy: Mainframe to PC Evolution



Vertically integrated  
Closed, proprietary  
Slow innovation  
Small industry




Horizontal  
Open interfaces  
Rapid innovation  
Huge industry



\* Slide courtesy: N. McKeown


# Traffic Engineering



Q: What if network operator wants u-to-z traffic to flow along uvwz, and x-to-z traffic to flow xwyz?

A: Need to define link weights so traffic routing algorithm computes routes accordingly (or need a new routing algorithm)!


# Traffic Engineering



Q: What if network operator wants u-to-z traffic to flow along uvwz and uxyz (load balancing)?

A: Can't do it (or need a new routing algorithm)

# Traffic Engineering



Q: What if w wants route blue and red traffic differently?

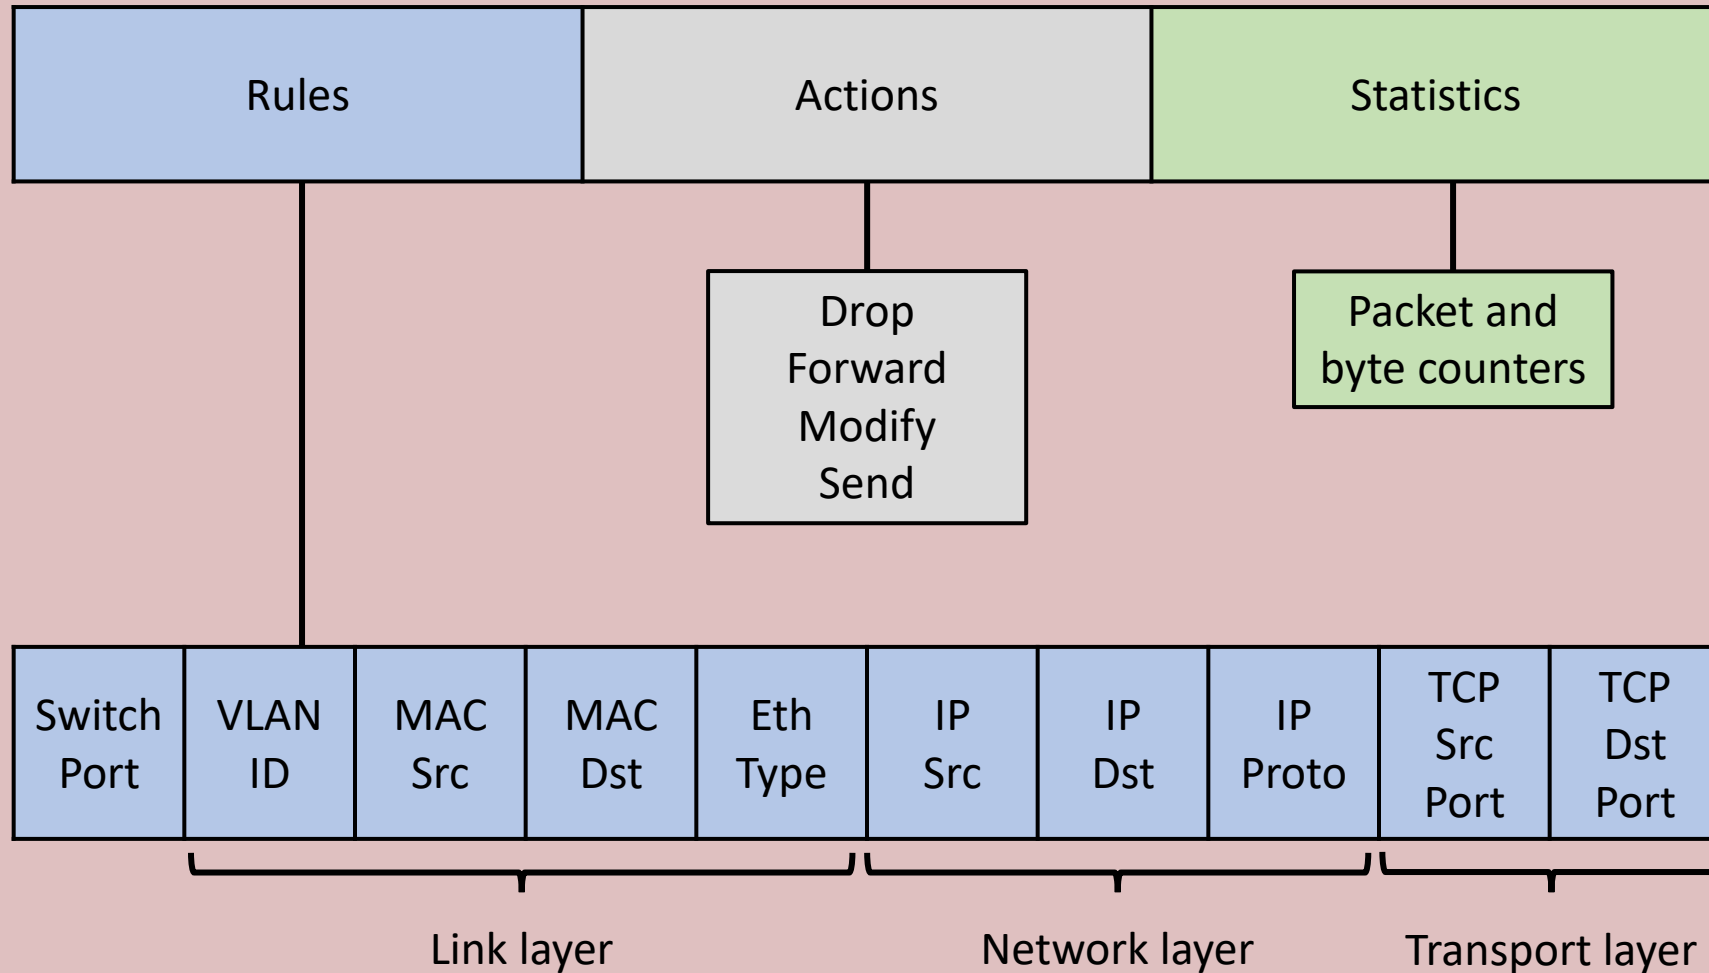
A: Can't do it (with destination based forwarding, and LS, DV routing)

# Review: Generalized Forwarding

- Logically centralized controller computes and distributes **flow tables**
- More than just forwarding packet to destination (**router**)
- Following actions are also supported
  - **Drop** packet (firewall)
  - **Modify** packet (NAT)
  - **Send** packet to SDN controller



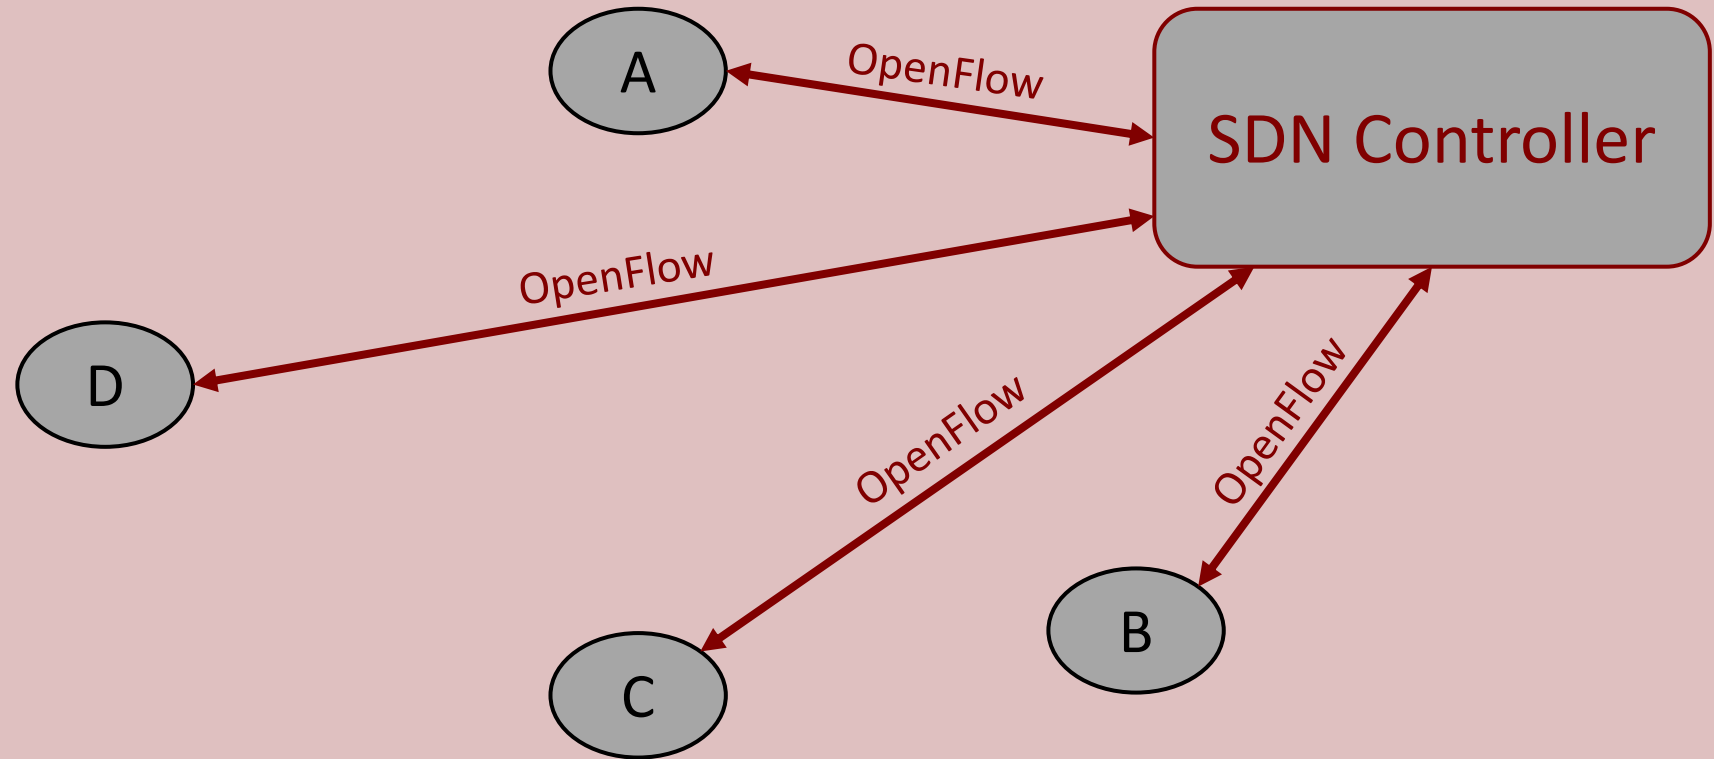
# Review: Flow Table



# Software-Defined Networking (SDN)

- Logically centralized **SDN controller**
- **OpenFlow**: protocol between SDN controller and routers; flow table format (**mix and match hardware vendors**)
- Routers/switches (same hardware) perform **generalized forwarding** based on flow tables from controller

# SDN



**Routers:** report statistics and topology updates

**SDN Controller:** computes and distributes flow tables

# Thank You!